

Optimal WTI—FX Hedge Ratios Under a Fixed Budget

Jihong Park

Pusan National University

adrianus026@gmail.com

July 2026

Abstract

A Korean crude-oil importer facing correlated WTI and USD/KRW exposure must decide, before any option is bought, what fraction of each exposure to hedge out of a fixed spending authority, and in which instrument. This paper poses the decision as a convex program and answers four questions, of which the instrument question carries the central result. A knock-out barrier option's premium discount exists because the option can die, and the scenario the hedge is sized against is the scenario in which it is most likely to die: re-simulated from the stress state with jump-adapted exact first-passage sampling, the contract actually held is dead on roughly half the paths that reach the scenario it was bought against. A mixed vanilla/knock-out program with the survival-adjusted ledger as its budget constraint prices this: it switches instruments bang-bang at a closed-form indifference mortality of 4.24%, an order of magnitude below the measurement, and selects the all-vanilla book at no cost in residual volatility. Aligning the two instruments' strikes strengthens the verdict to dominance: the same-strike knock-out costs more than the vanilla under this paper's pricing conventions, so the apparent discount prices the strike gap and not the barrier. Around this result the allocation geometry is mapped in closed form. The extreme 97/3 WTI/FX split follows from an 18-to-1 variance ratio with near-zero correlation, invariant across premium regimes; cost minimization selects a different, worse-hedged corner than risk minimization; and whether the assumed KRW 45bn authority binds at all is regime- and specification-conditional, a vertex under the European specification and slack under the American, undone by any one of four independent, individually plausible respecifications, so the budget-geometry answer is reported as the conditional result it is. The reported split is likewise conditional on a calm-sample correlation estimate and a one-notional coverage cap, both traced rather than assumed.

Keywords: corporate hedging; hedge ratio; convex optimization; knock-out barrier; quanto; commodity and currency risk; budget constraint

JEL classification: G32; G13; Q40; C61

Contents

1 Introduction

3

2	Literature Review	4
3	The Static Allocation Problem	5
3.1	Decision Variables and Risk Objective	5
3.2	Cost Functions	6
3.3	Is the premium split identified? A degeneracy audit	7
3.4	Constraint Set	8
4	Calibration	9
5	Numerical Method and Verification	11
6	Numerical Results	12
6.1	Question (a): Where the Budget Binds, a Vertex in One Regime and Slack in the Other	12
6.2	Question (b): The Mechanism Behind the Extreme w_1/w_2 Asymmetry	13
6.3	Question (c): Cost Minimization Is a Different Program with a Different Answer	15
6.4	The Global Geometry	15
6.5	Parameter Robustness	16
6.6	How much of the answer is the allocation constraint? A κ -sweep	18
6.7	A co-struck, co-termed respecification	19
7	The Knock-Out Survival Haircut	20
8	The Mixed Vanilla/Knock-Out Program	22
9	Discussion	25
10	Practical Implications	27
11	Limitations	28
12	Conclusion	29

1 Introduction

Before a hedging desk ever debates *how* to run an option hedge (which pricing model, which rebalancing rule, which Greeks) the treasurer has to answer a cruder and more consequential question: out of a fixed hedging budget, *what fraction of each exposure should be covered at all?* For a Korean crude-oil importer the exposures are two: the WTI price of the physical barrels it must buy every month, and the USD/KRW exchange rate at which the resulting dollar liability is converted into won. In the configuration studied here the physical requirement is 2,000,000 barrels per month, the equivalent dollar need is \$157,880,000 per month, and the monthly option budget is set at an assumed KRW 45,000,000,000, a figure whose provenance and consequences Section 4 states explicitly. The decision variables are two scalars: $w_1 \in [0, 1]$, the fraction of the WTI exposure covered by WTI call options, and $w_2 \in [0, 1]$, the fraction of the USD need covered by USD/KRW call options. This is a *static* allocation made once, before the hedge is put on; nothing in it depends on how the position is subsequently managed.

The economics of the trade-off are stark on both sides. Hedging is expensive: covering the full oil leg with Black-76-priced calls costs KRW 41.3bn (nearly the entire budget), and covering the full FX leg costs KRW 13.8bn with the Garman—Kohlhagen premium or KRW 4.2bn with the American structure’s Shapley-attributed currency premium. Not hedging is also expensive: under the stress scenario against which the program is sized (WTI at 113 USD/bbl, USD/KRW at 1550), the uncovered oil exposure alone loses KRW 105.6bn if nothing is hedged. The allocation problem is the collision of these two ledgers under one spending cap.

Four questions organize the paper. The first three map the allocation: (a) **is the budget constraint actually binding**, a boundary point held back by money or an interior risk optimum that merely sits near the cap; (b) **what mechanism produces the extreme w_1/w_2 asymmetry** the solutions exhibit, an artifact of the optimizer or a structural consequence of the volatility and correlation inputs; and (c) **would minimizing cost directly give the same allocation** as minimizing risk subject to the budget, and if not, which criterion should govern. The fourth is the paper’s central question, and it concerns the instrument rather than the allocation: (d) **does the knock-out structure’s premium discount survive its own mortality?** A knock-out option’s premium is low precisely because the option can die, and the scenario in which protection is most needed is also the scenario in which death is most likely. The paper re-simulates the calibrated jump-diffusion from the stress state itself to measure that mortality directly, rather than quoting unconditional knock-out statistics, and then prices it into the allocation through a mixed vanilla/knock-out program in which the survival-adjusted ledger is the budget constraint and the optimizer chooses the instrument as well as the coverage.

Both allocation programs are convex, so their global solutions can be characterized in closed form and verified numerically. The answers, previewed. (d) is the headline: the measured stress mortality is 0.49 for the contract actually held, eleven times the closed-form indifference point at which the premium discount stops paying for its own mortality, so the mixed program selects the all-vanilla book, and once the two instruments are struck alike the knock-out carries no discount at all (Section 6.7). (b) the asymmetry is structural and can be written in one closed-form line, with the FX leg left 96.6% open by the variance arithmetic itself. (c) cost minimization is a different program with a different (corner) answer, and the two criteria coincide only when

the cost program is re-anchored by the risk cap it discards. (a) is the conditional answer: the combined premium-plus-stress ledger binds at a vertex in the European regime, shallowly, in the sense that unlimited budget buys almost no further variance reduction, and is slack in the American, and the vertex is undone by any of four independent respecifications (Sections 4, 6.6, 6.7), so the budget-geometry result is reported with its conditions attached.

2 Literature Review

The objective function is the classical minimum-variance portfolio criterion of Markowitz (1952), specialized to two residual exposures: the quantity minimized is the standard deviation of a two-asset portfolio whose weights are the *unhedged* fractions $(1 - w_1)$ and $(1 - w_2)$. The closed-form line-restricted solution used in Section 6.2 to explain the allocation asymmetry is the textbook two-asset global-minimum-variance weight.

The premium inputs on the European side come from two closed-form option models. Black (1976) prices the WTI leg as a call on the commodity forward; Garman and Kohlhagen (1983) price the FX leg as a currency call with separate domestic and foreign interest rates. On the American side, premiums come from a Longstaff—Schwartz (2001) least-squares Monte Carlo valuation of a knock-out structure under jump-diffusion dynamics in the spirit of Merton (1976); the diffusive volatility 0.3242 that valuation consumes is what a threshold-classification jump/diffusion decomposition leaves after stripping the jump variance out of the raw historical volatility 0.3946, and it is an input to the pricing engine and not to the risk objective, which takes the raw figure because the uncovered barrels bear the jumps as well as the diffusion. The joint quanto premium is never assembled from a closed-form correlation-drift adjustment of the kind derived by Reiner (1992); it is instead recovered directly from the correlated joint simulation of Section 4, which reproduces the same coupling automatically once S_1 and S_2 are drawn from correlated shocks. Because the American structure prices both legs jointly, the joint premium is attributed to the individual legs by Shapley (1953) value, producing per-leg premium constants that enter the cost function linearly. The premiums are priced at inception, at today's spot, deliberately so: the budget constraint governs cash paid *today*, and re-pricing the premium at the stress spot would charge the stress scenario twice, once in the premium and once in the unhedged-loss terms that already price it.

The economic question (how much of each of two correlated exposures to cover when hedging is costly) descends from the cross-hedging literature of Anderson and Danthine (1981) and, for the currency dimension of internationally sourced commitments, Adler and Dumas (1984). The present setting differs from the classical futures cross-hedge in one decisive respect: hedging here is done with options bought under a hard spending authority, so the operative trade-off runs between variance and an explicit cap on a combined premium-plus-stress ledger, in place of the usual variance/expected-return frontier, which is precisely what turns a one-line minimum-variance formula into a constrained program worth studying. The unconstrained special case of that formula is nonetheless the classical minimum-variance hedge ratio of Ederington (1979); the budget cap of Section 3.3 is what separates the present, corner- and vertex-seeking program from a direct application of that one-line result. The closest instrument-level antecedent is Ahn,

Boudoukh, Richardson, and Whitelaw (1999), who optimize a put-option hedge under an explicit premium constraint and characterize the strike that buys the most value-at-risk reduction per dollar of cost; the present program shares their premium-constrained logic but allocates one budget across two correlated exposures with the contracts held fixed, so the decision variable is coverage rather than contract design. Brown and Toft (2002) solve the complementary problem, deriving the firm’s optimal customized payoff under joint price and quantity risk; here the payoffs are given market instruments, which is what makes the question an allocation rather than a design problem.

The premium budget is taken here as exogenous, following the corporate hedging literature that supplies its rationale. Rampini, Sufi, and Viswanathan (2014) derive hedging bounded by collateral and financing capacity rather than chosen at an unconstrained optimum, which is the mechanism a spending authority proxies for; Campello, Lin, Ma, and Zou (2011) document the financing frictions empirically, and Pérez-González and Yun (2013) and Gilje and Taillard (2017) supply the identification-based evidence that hedging affects value. Froot, Scharfstein, and Stein (1993) model this kind of institutional friction, though they endogenise the hedging decision rather than take a spending authority as given, so they motivate the constraint’s existence without licensing its exogeneity: a firm hedges to protect its internal financing capacity for investment and not to maximize variance reduction per se, which is consistent with hedging being bounded by an approved cash outlay and not pursued to a first-order risk optimum. Stulz (1996), arguing from comparative information advantage rather than from variance arithmetic, develops the practitioner-facing corollary (selective, partial hedging can be optimal even when full coverage is affordable), which is the same qualitative conclusion Section 6.2’s closed-form asymmetry reaches by a different, purely variance-arithmetic route. Jin and Jorion (2006), whose event study addresses firm value and the hedging premium rather than coverage levels, supply the closest empirical anchor to the present institutional setting, documenting derivative-based hedging behavior among U.S. oil and gas producers; the importer studied here sits on the opposite side of the same commodity risk, and the partial coverage that Sections 6—8 derive from first principles for this specific instrument set is consistent with the incomplete hedging their sample exhibits, though they document no budget-bounded pattern.

3 The Static Allocation Problem

3.1 Decision Variables and Risk Objective

Let σ_1 denote the annualized volatility of WTI log returns, σ_2 that of USD/KRW log returns, and ρ their correlation. If fractions w_1 and w_2 of the two exposures are hedged, the unhedged remainders $(1 - w_1)$ and $(1 - w_2)$ carry the residual risk

$$\sigma_{res}(w_1, w_2) = \sqrt{(1 - w_1)^2 \sigma_1^2 + (1 - w_2)^2 \sigma_2^2 + 2(1 - w_1)(1 - w_2) \sigma_1 \sigma_2 \rho}. \quad (1)$$

Both regimes share every risk input. σ_{res} in (1) measures the volatility of the exposure left uncovered, and that exposure does not know which model priced its hedge, so $\sigma_1 = 0.3946$, $\sigma_2 = 0.0926$ and $\rho = 0.0876$ enter it whichever instrument is bought. The jump-stripped

diffusive figure $\sigma_1^{diff} = 0.3242$ appears only inside the American Monte Carlo engine, which restores the jump component separately; taking it into the risk objective would credit the uncovered barrels with a jump immunity they do not have.

3.2 Cost Functions

Total cost is the sum of two ledgers: the premium paid today for the covered fraction, and the stress-scenario loss suffered by the uncovered fraction. Both regimes share the stress ledger and differ only in the premium ledger. With $Q_{oil} = 2,000,000$ bbl, $Q_{USD} = \$157,880,000$, spot prices $S_{WTI} = 78.94$ and $S_{KRW} = 1540.64$, stress levels $S_{WTI}^{st} = 113$ and $S_{KRW}^{st} = 1550$, funding rate $r_w = 0.07$ (WACC), and maturities $T_1 = 0.833$, $T_2 = 0.5$ years:

$$\begin{aligned} C^{EU}(w_1, w_2) = & w_1 Q_{oil} P_{B76} S_{KRW} (1 + r_w T_1) + w_2 Q_{USD} P_{GK} (1 + r_w T_2) \\ & + (1 - w_1) Q_{oil} \max(0, S_{WTI}^{st} - S_{WTI}) S_{KRW}^{st} \\ & + (1 - w_2) Q_{USD} \max(0, S_{KRW}^{st} - S_{KRW}), \end{aligned} \quad (2)$$

$$\begin{aligned} C^{AM}(w_1, w_2) = & w_1 Q_{oil} P_{WTI}^{Sh} (1 + r_w T_1) + w_2 Q_{oil} P_{FX}^{Sh} (1 + r_w T_2) \\ & + (1 - w_1) Q_{oil} \max(0, S_{WTI}^{st} - S_{WTI}) S_{KRW}^{st} \\ & + (1 - w_2) Q_{USD} \max(0, S_{KRW}^{st} - S_{KRW}), \end{aligned} \quad (3)$$

with $P_{B76} = 12.6524$ USD/bbl (Black-76 WTI call), $P_{GK} = 84.667$ KRW per USD (Garman—Kohlhagen FX call), and $P_{WTI}^{Sh} = 15,093.75$, $P_{FX}^{Sh} = 2,038.72$ KRW/bbl. The FX leg is scaled by Q_{USD} in (2) because P_{GK} is quoted per dollar, and by Q_{oil} in (3) because the Shapley attribution splits a per-barrel joint premium; the notional identity (7), $Q_{oil} S_{WTI} = Q_{USD}$, is what makes $w_2 = 1$ the same dollar of coverage in both regimes, so the two cash ledgers are comparable leg by leg (LSMC Shapley attribution of the joint per-barrel premium 17,132.47 KRW/bbl, so both shares scale on the barrel position). The split is the Shapley value of the two-player game whose characteristic function is the joint quanto premium, $v(\{i\})$ being the structure repriced with the other leg's volatility set to zero and $v(\{1, 2\})$ the joint premium. Section 3.3 shows that for this particular payoff that characteristic function is degenerate, that the resulting split is therefore a reporting convention rather than an identified quantity, and exactly which of the paper's results do and do not depend on it.

The attribution is a *cost-reporting* convention and not a market price, and Section 8 turns on exactly that distinction. Sections 6 and 7 price the American structure as one indivisible instrument and use the Shapley shares to report how its single premium is charged to the two coverage ratios; no leg is assumed separately tradable there. Section 8 does assume separate tradability, because the mixed program buys a standalone knock-out leg alongside a standalone vanilla leg, and for that purpose the Shapley share is the wrong number: an independently priced single-asset KO call comes to KRW 17,377.11/bbl, 15.1% above the attributed share. The two sections therefore use different prices because they are pricing different things, a jointly-held structure and a separable pair of legs. Both regimes carry their cash premiums to the decision horizon on the same simple-interest factor $(1 + r_w T)$, so the two ledgers are directly comparable

leg by leg. Both functions are affine, and collecting terms gives the marginal cost of each ratio:

$$\frac{\partial C^{EU}}{\partial(w_1, w_2)} = (-64.327, +12.357) \text{ bn KRW}, \quad \frac{\partial C^{AM}}{\partial(w_1, w_2)} = (-73.638, +2.745) \text{ bn KRW}. \quad (4)$$

The signs carry the economics. Raising the WTI ratio *reduces* total cost in both regimes: the avoided stress loss (KRW 105.59bn at full exposure) dwarfs the full-cover premium (KRW 41.26bn European, KRW 31.95bn American). Raising the FX ratio *increases* it: in both regimes the FX premium exceeds the small FX stress gap (KRW 1.48bn at full exposure), by KRW 12.36bn per unit of w_2 in the European regime and KRW 2.75bn in the American, where the barrier structure prices the currency leg far more cheaply. Equation (4) drives everything in Section 6.3; the American column's FX entry is priced off the Shapley attribution, and Section 3.3 audits how much of it is convention.

3.3 Is the premium split identified? A degeneracy audit

The Shapley construction prices $v(\{1\})$ by freezing the currency volatility and $v(\{2\})$ by freezing the WTI dynamics. For this payoff the first operation does nothing. The companion paper's factorization result (Park 2026b, Proposition 1) shows the structure's value is $S_2(0)e^{-r_{US}T} \mathbb{E}^{Q_{US}}[X(T)]$, in which neither σ_2 nor ρ appears, so

$$v(\{1\}) = v(\{1, 2\}) \quad \text{exactly, and hence} \quad \phi_{FX} = \frac{1}{2} v(\{2\}), \quad (5)$$

half the value of the structure with WTI frozen: a deterministic intrinsic-value quantity, generated by the drift of a volatility-free WTI path, containing no FX optionality at all. The pricing engine confirms the degeneracy directly: at 100,000 paths over two seeds, $v(\{1, 2\}) - v(\{1\})$ is 77–97 KRW per unit against a Monte Carlo standard error of 63 and a premium near 17,100, which is the size of the engine's documented drift-tilt inconsistency and not of any FX risk price, and ϕ_{FX} sits within 50 KRW of $\frac{1}{2}v(\{2\})$. The joint premium simply has no FX-volatility component to attribute, and the Shapley axioms cannot conjure one; what they return is a bookkeeping share.

The split is therefore one convention among several, and the conventions disagree materially. An exposure-notional attribution, which charges each leg its own currency of risk using the companion paper's structural deltas (WTI leg $\delta_1 S_1 S_2 \approx 64,968$ KRW/bbl of exposure against the option's own currency exposure $S_2 \partial V / \partial S_2 = V \approx 17,132$), puts the FX share at 20.9%, or 3,575 KRW/bbl; a variance-contribution attribution puts it at 0.9%, or 156 KRW/bbl; Shapley reports 11.9%, or 2,038.72.

What survives the ambiguity and what does not. Every attribution sums to the joint premium, so any quantity evaluated where the two ratios move together is convention-invariant: the full-coverage total (KRW 36.17bn at $w_1 = w_2 = 1$) and, decisively for question (a), the American regime's budget slack. Re-costing the line optimum (0.9660, 0.0340) under each convention above gives totals from KRW 32.99bn to KRW 40.05bn, the worst case being the one that charges the entire premium to the WTI leg, and all of them sit inside the KRW 45bn authority: the regime contrast of Section 6.1 does not depend on the split. What does depend

on it is every *marginal* statement. $\partial C^{AM}/\partial w_2$ ranges from -1.15bn per unit of w_2 under variance attribution through $+2.75\text{bn}$ under Shapley to $+5.92\text{bn}$ under exposure attribution, the sign included, so the reading of (4) as “the barrier structure prices the currency leg cheaply” is a property of the convention and not of the structure. Section 7’s Shapley-priced threshold (12) inherits the same caveat, which is why Section 8 re-derives its thresholds from standalone, separately-traded instrument prices, and why every decision this paper actually recommends, the instrument switch and the mortality comparison, is made on those prices rather than on the attribution.

3.4 Constraint Set

The allocation program imposes exactly four constraints:

$$0 \leq w_1 \leq 1, \quad 0 \leq w_2 \leq 1, \quad w_1 + w_2 \leq 1, \quad C(w_1, w_2) \leq B = 45,000,000,000 \text{ KRW}, \quad (6)$$

with C equal to (2) or (3).

What B caps. One point should be fixed before any result is read. $C(w_1, w_2)$ in (6) is not cash paid at inception. It is the sum of the option premium, which is cash, and the stress-scenario loss on the unhedged fraction, which is a contingent outcome under one scenario. At the European risk-minimising vertex the two components are KRW 40.45bn of premium and KRW 4.55bn of stress loss, summing to the KRW 45bn cap. Under a cap on *premium alone* the same allocation spends 40.45bn, the unconstrained line optimum (0.9660, 0.0340) spends 40.33bn, and the cap binds in neither regime.

The vertex result of Section 6.1 is therefore a statement about a combined spend-plus-stress ledger and not about a cash authority, and the reader should hold that distinction throughout. A treasury whose 45bn authority governs cash outlay only would find both regimes interior and the allocation pinned by (9) alone, which is the American answer (0.9660, 0.0340) in both cases. The paper reports the combined-ledger program because that is what makes the two regimes’ stress exposure comparable on one axis, but the “budget binds” language of question (a) applies to that ledger and not to a premium budget.

The allocation constraint $w_1 + w_2 \leq 1$ deserves more than a passing note, because the paper’s central asymmetry result is derived on the line it defines. Its economic content is that the two legs are denominated against *the same* exposure. At the spot rate of Table 1 the monthly barrel requirement and the monthly dollar requirement are the same quantity of money,

$$Q_{oil} \cdot S_{WTI} = 2,000,000 \times 78.94 = \$157,880,000 = Q_{USD}, \quad (7)$$

so a barrel hedged and a dollar hedged cover the identical underlying cash outflow, once through its price and once through its currency of settlement. A treasury that authorizes hedging of “the monthly crude bill” and treats w_1 and w_2 as two ways of covering that one bill is imposing precisely $w_1 + w_2 \leq 1$. Read this way the constraint is a statement about what is being hedged rather than an arbitrary cap.

The reading is not forced, however, and the results are sensitive to it. A firm that regarded the price exposure and the currency exposure as two separate books, each hedgeable up to its own notional, would impose $w_1 + w_2 \leq \kappa$ with $\kappa = 2$, and Section 6.6 shows what that does to the answer. The constraint is therefore reported here as a modelling choice with an economic rationale, and its consequences are traced rather than assumed away.

Together with the budget cap these constraints carve the feasible polygon whose geometry Section 6 maps completely: a sliver satisfying $w_1 \geq 0.9648 + 0.1921 w_2$ in the European regime, and a far roomier, nearly horizontal band satisfying $w_1 \geq 0.8434 + 0.0373 w_2$ in the American regime.

4 Calibration

Table 1 lists every input used in the paper. The underlying series are daily WTI spot (USD/bbl) and USD/KRW, covering 28 June 2021 to 22 June 2026, 1,301 paired closing observations retained after aligning the two calendars and dropping non-common trading days. Differencing gives 1,300 paired daily log returns, of which 1,299 survive the jump-classification screen of Section 7 and constitute the estimation sample used throughout. The risk inputs are computed from those 1,299 paired daily log returns: $\sigma_1 = 0.3946$ (WTI, annualized) and $\sigma_2 = 0.0926$ (USD/KRW), with correlation $\rho = 0.0876$. The American regime's $\sigma_1 = 0.3242$ is the diffusive volatility from the threshold-classification jump/diffusion decomposition of the same series: jump days are identified by a k -sigma rule, the two jump regimes are moment-matched on the classified days, and the diffusive volatility is re-estimated from the remainder, so the observed variance is split rather than duplicated (the construction is documented in the companion paper). The two closed-form legs are struck 5% in the money, at $0.95 S_{WTI} = 74.99$ USD/bbl and $0.95 S_{KRW} = 1,463.61$ KRW/USD; the American structure is struck at the money, $K = S_{WTI} = 78.94$, with knock-out barriers at 120 and 50 USD/bbl. The two regimes therefore price different strikes, and the two legs carry different maturities ($T_1 = 0.833$ on the WTI leg against $T_2 = 0.5$ on the FX leg), both specification choices inherited from the production configuration rather than derived, so the leg-by-leg cash comparisons of (4) embed both mismatches; Section 6.7 removes both mismatches and re-solves the program: the premium levels move, the closed-form asymmetry (9), which reads no contract parameter at all, does not, and the instrument comparison of Section 8 sharpens.

Why the premium is large, and what that costs question (a). The European premium of $P_{B76} = \$12.65/\text{bbl}$ is 16.6% of the monthly crude bill, which is high against corporate practice. The pricing is not at fault: recomputing Black-76 at the Table 1 inputs gives $d_1 = 0.3225$, $d_2 = -0.0376$ and $\$12.6524$, reproducing the engine exactly. The level comes from three contract choices, each of which is defensible on its own and which multiply.

The strike is 5% in the money, so $\$3.95$ of the premium, 31% of it, is intrinsic value paid up front rather than a price for transferring risk; the time value is $\$8.71$. The maturity is 0.833 years against a *monthly* exposure, and premium scales roughly with $\sqrt{T}$. And $\sigma_1 = 39.46\%$ is realized volatility over a window containing the 2022 spike, whereas an option would be priced

off implied volatility, typically 30–35% for WTI. Repricing under market-convention choices, 5% out of the money at one month and 33% volatility, gives \$1.49/bbl, and full WTI cover then costs KRW 4.86bn against the same authority.

That last number is what question (a) rests on. Under the paper’s specification, full European WTI cover costs KRW 41.26bn and the KRW 45bn authority binds; under the market-convention specification it costs KRW 4.86bn and the authority is slack in both regimes, so the answer to question (a) becomes “neither binds” and the regime contrast disappears. Question (a) is therefore conditional on the contract specification in the same way it is conditional on B (8) and on κ (Section 6.6), and the three conditions are independent routes to the same reversal.

The specification is retained rather than changed because the paper’s remaining results do not depend on it. The closed-form asymmetry (9) reads only σ_1/σ_2 and ρ ; the cost-minimization contrast of question (c) compares two corners of the same ledger; and the instrument-mix indifference point $\bar{p} = (P_V - P_K)/UL_{WTI}$ of Section 8 is a ratio of premium *differences*, so re-specifying both legs moves its level without changing its sign or the direction of the switch. What a market-convention specification would remove is the budget geometry of Section 6.1, and only that.

Why options rather than futures. A minimum-variance hedger would reach for futures, which carry no premium and remove most of the variance, or for a zero-cost collar. The premium in (2) is the price of asymmetry: an importer gains when crude falls, and a futures hedge surrenders that gain along with the loss. The budget constraint studied here exists because the firm has chosen to keep the downside participation that a linear hedge would give away, and the paper takes that choice as given rather than deriving it. A firm indifferent to the asymmetry has no budget problem to solve, and should hedge linearly.

The spending authority B . The KRW 45bn figure is an assumption of this paper and not an observed corporate disclosure, and because question (a) turns on it the assumption is stated as a rule rather than as a constant. The monthly crude bill at the calibration spot is $Q_{USD} \cdot S_{KRW} = \text{KRW } 243.2\text{bn}$, so B is 18.5% of the monthly procurement outlay, and at the European vertex the premium component alone is 16.6% of it. That is a large authority by the standards of corporate hedging practice, where option budgets are more commonly a low single-digit percentage of the exposure being covered, and it is chosen here because a smaller one makes the European program infeasible against the vanilla premiums of Table 1 rather than because any firm is known to operate at that level. A reader who regards 18.5% as implausible should read the European regime as the illustrative one and the American regime, which spends KRW 36.02bn or 14.8% of the bill, as the operative case.

The consequences are sharp enough to quantify, and Section 11 records them. The European vertex of Section 6.1 exists only for

$$B \in [42.74, 45.35] \text{ bn KRW}, \quad (8)$$

an interval 6.1% wide. Below KRW 42.74bn the cost boundary no longer reaches the allocation

line and the program is infeasible at full WTI cover; above KRW 45.35bn the budget-free line solution becomes affordable, the cap goes slack, and the answer to question (a) reverses from “binds at a vertex” to “does not bind”. The assumed KRW 45bn sits 0.77% below that upper endpoint. This is a narrower margin than any other assumption in the paper carries, and question (a) should be read as conditional on it.

The stress scenario (WTI at 113 USD/bbl and USD/KRW at 1550) prices the unhedged fraction of each leg in (2)—(3).

Quantity	Symbol	Value
WTI spot (USD/bbl)	S_{WTI}	78.94
USD/KRW spot	S_{KRW}	1,540.64
Monthly oil need (bbl)	Q_{oil}	2,000,000
Monthly USD need	Q_{USD}	157,880,000
WACC (funding rate)	r_w	0.07
Budget (KRW)	B	45,000,000,000
European WTI call strike (USD/bbl)	$K_1 = 0.95 S_{WTI}$	74.99
European FX call strike (KRW/USD)	$K_2 = 0.95 S_{KRW}$	1,463.61
American call strike (USD/bbl)	$K = S_{WTI}$	78.94
Knock-out barriers (USD/bbl)	U, L	120, 50
WTI option maturity (yr)	T_1	0.833
FX option maturity (yr)	T_2	0.5
Stress WTI (USD/bbl)	S_{WTI}^{st}	113
Stress USD/KRW	S_{KRW}^{st}	1,550
WTI volatility, risk objective	σ_1	0.3946
WTI diffusive vol, pricing engine	σ_1^{diff}	0.3242
USD/KRW vol	σ_2	0.0926
WTI—FX correlation	ρ	0.0876
USD risk-free rate (pricing)	r_{US}	0.040
KRW risk-free rate (pricing)	r_{KRW}	0.035
Jump intensity (symmetrized)	λ	6.7846 /yr
Mean jump size (log)	θ_J	−0.0299
Jump volatility (log)	δ_J	0.08443
Black-76 WTI call (USD/bbl)	P_{B76}	12.6524
Garman—Kohlhagen FX call (KRW/USD)	P_{GK}	84.667
LSMC Shapley WTI premium (KRW/bbl)	P_{WTI}^{Sh}	15,093.75
LSMC Shapley FX premium (KRW/bbl)	P_{FX}^{Sh}	2,038.72

Table 1: Model inputs. The risk and cost blocks are estimated from the price history described in Section 4; the jump block and the two pricing rates enter only the American premium and the knock-out simulation of Sections 7–8. Every downstream number in this paper is computed from these values.

5 Numerical Method and Verification

Both allocation programs are convex: σ_{res} in (1) is the Euclidean norm of an affine map of (w_1, w_2) , hence convex; the cost functions (2)—(3) are affine; and every constraint in (6) is affine. Any local minimum is therefore the global minimum, and the solution can be established analytically and not only searched for.

Two complementary numerical approaches confirm the global optima. A deterministic grid search (801 points per axis, twice refined) covers the feasible region exhaustively. In parallel, a

sequential least-squares quadratic programming routine (SLSQP, 73 multi-starts) recovers the boundary solutions alongside closed-form analytic derivations from the affine constraint algebra. Across all evaluated scenarios the two solvers and the analytic bounds locate the same vertices.

The convex nature of the programs also permits a consistency check, exploited throughout Section 6. It should not be read as independent verification: minimizing cost subject to $\sigma_{res} \leq \sigma_{res}^*$, where σ_{res}^* is the achieved minimum and σ_{res} is strictly decreasing in both ratios over the polygon, admits a unique feasible point, so any working solver must return it. The check catches solver failure and cannot fail for any other reason. The cost minimum taken subject to a risk cap equal to the achieved minimum risk returns the risk-minimizing point itself. Both regimes satisfy this duality constraint, consistent with the derived allocations being the constrained optima.

Two further computations feed Sections 7 and 8. The scenario-conditional knock-out probability of Section 7 is measured by direct simulation of the calibrated WTI jump-diffusion (diffusive volatility 0.3242, jump intensity $\lambda = 6.7846$ per year, jump mean -0.0299 , jump volatility 0.08443, physical drift 0.139, Merton compensator in the drift) started *at the stress spot* and monitored against the 120/50 barriers over the option’s 0.833-year life with jump-adapted exact first-passage sampling: Poisson jump times are drawn exactly within each daily step, the Brownian-bridge crossing probability is applied on each jump-free sub-interval, and the barrier is checked directly at each jump instant, so the reported touch probability carries no time-discretization bias (200,000 paths; coarser monitoring understates this probability materially, since a barrier that is touched and re-crossed inside a recording gap leaves no trace in the recorded values). The mixed-instrument program of Section 8 exploits an exact reduction: its survival-adjusted ledger is linear in the vanilla/KO split at any fixed total coverage, so the optimal split is always at an endpoint (all-vanilla or all-KO), and the two-instrument program collapses to the lower envelope of two two-variable convex programs solved with the same machinery.

6 Numerical Results

Table 2 collects the six solved programs together with the two budget-relaxed benchmarks; the subsections interpret them.

6.1 Question (a): Where the Budget Binds, a Vertex in One Regime and Slack in the Other

Three computations characterize the budget’s role.

First, the active set. At the European risk minimum the allocation constraint and the budget constraint hold with equality *simultaneously*: $w_1^* + w_2^* = 1$ and $C(w_1^*, w_2^*) = B$. The optimum is the vertex where the budget boundary crosses the allocation line, visible in the zoomed panels of Figure 3, where the iso-risk contours close in on a point pinched between the solid budget boundary and the dashed allocation line. The American regime has a different active set: only $w_1^* + w_2^* = 1$ holds, at a total cost of KRW 36.02bn, KRW 8.93bn inside the cap, so the budget boundary runs well below the solution and never touches it.

Regime	Program	w_1	w_2	Total cost (KRW bn)	σ_{res}
European	risk min. (budgeted)	0.9705	0.0295	45.000	0.0916
European	risk min. (no budget)	0.9660	0.0340	45.346	0.0916
European	cost min.	1.0000	0.0000	42.737	0.0926
European	cost min. $\leq$ risk cap	0.9705	0.0295	45.000	0.0916
American	risk min. (budgeted)	0.9660	0.0340	36.024	0.0916
American	risk min. (no budget)	0.9660	0.0340	36.024	0.0916
American	cost min.	1.0000	0.0000	33.425	0.0926
American	cost min. $\leq$ risk cap	0.9660	0.0340	36.024	0.0916

Table 2: Solved allocation programs. The σ_{res} column is instrument-blind: it credits w_1 with full protection whether the WTI leg is a vanilla call or a knock-out. Section 8 shows that crediting the American leg with survival-adjusted coverage instead raises its residual volatility from 0.0916 to 0.372 at the measured stress mortality, so the American rows of this table should be read as the program’s own objective value and not as the risk the American book actually carries. In the European regime the budgeted risk minimum is a vertex solution: the allocation constraint $w_1 + w_2 \leq 1$ and the budget constraint hold with equality simultaneously. In the American regime the budget is slack at the optimum, so the allocation constraint binds alone and the budgeted and unbudgeted minima coincide. The risk-capped cost minima reproduce the risk minima (the duality audit of Section 5).

Second, the relaxation test. Deleting the budget constraint moves the European optimum from (0.9705, 0.0295) to (0.9660, 0.0340), which lowers σ_{res} only slightly; funding that relaxed optimum would cost KRW 45.35bn, 0.77% over budget and so unaffordable. The cap therefore genuinely excludes the unconstrained European optimum, and the ledger constraint is binding there in the exact Karush—Kuhn—Tucker sense. The American optimum does not move at all: its unconstrained minimum (0.9660, 0.0340) already costs only KRW 36.02bn, some 20% inside the cap, so the constraint is inactive and its multiplier is zero. The regimes differ because the American ledger prices the currency leg at KRW 2.75bn per unit of w_2 against the vanilla leg’s KRW 12.36bn, so the FX coverage the risk objective wants is affordable in one regime and not in the other; the marginal price is convention-dependent (Section 3.3), the slack itself is not.

Third, the price of the constraint. Figure 1 traces the optimal σ_{res} as the budget sweeps from KRW 33bn to 58bn. Each curve falls, then flattens at the relaxed-optimum budget just computed (KRW 45.35bn European, KRW 36.02bn American), beyond which money buys nothing. The KRW 45bn cap sits on the falling branch of the European curve and on the flat branch of the American one. Finite-difference shadow prices at $B = 45$ bn are accordingly about one basis point of σ_{res} per additional billion KRW in the European regime and zero in the American. Where it binds the constraint is moreover *shallow*: even unlimited budget improves European residual volatility by 0.02%. The treasurer’s cap costs the firm almost no variance; conversely, requests for more hedging budget cannot be justified on variance grounds in this configuration.

6.2 Question (b): The Mechanism Behind the Extreme w_1/w_2 Asymmetry

The extreme tilt toward the WTI leg follows from three structural facts and not from any preference of the optimizer, each visible in closed form.

Fact 1: the WTI leg carries almost all the risk. $\sigma_1^2/\sigma_2^2 = 18.2$. At the no-hedge point (0, 0),




Figure 1: Budget relaxation sweep. The optimal residual volatility falls as the cap is raised and flattens at the unconstrained floor, reached at KRW 45.35bn (European) and KRW 36.02bn (American). The KRW 45bn cap (dashed) sits on the falling branch of the European curve, binding but close to the floor, and on the flat branch of the American curve, where it is slack.

where $\sigma_{res} = 0.4131$, the gradient of (1) is

$$\nabla \sigma_{res}(0, 0) = (-0.3846, -0.0285),$$

i.e. the first won of hedging is $13.5\times$ more productive on the WTI leg. A budget spent greedily on marginal risk reduction floods the WTI leg first.

Fact 2: the correlation is too small to couple the legs. The interaction term in (1) carries the factor $\sigma_1\sigma_2\rho = 0.0032$ (European), two orders of magnitude below σ_1^2 . With $\rho = 0.0876$, hedging WTI does almost nothing to the FX leg's marginal value and vice versa: the problem nearly decouples into two independent legs competing for one budget, and the leg with $18\times$ the variance wins almost all of it.

Fact 3: the allocation line pins the split in closed form. Because σ_{res} is strictly decreasing in both w_1 and w_2 (its gradient is negative throughout the unit square), any optimum must exhaust the binding envelope: the line $w_1 + w_2 = 1$, intersected with the budget. On that line, substituting $w_2 = 1 - w_1$ into (1) gives the classical two-asset minimum-variance problem in the unhedged weights $h_1 = 1 - w_1$, $h_2 = w_1$, whose solution is

$$w_1^{line} = \frac{\sigma_1^2 - \rho\sigma_1\sigma_2}{\sigma_1^2 + \sigma_2^2 - 2\rho\sigma_1\sigma_2} = 0.9660 \quad (\text{both regimes}). \quad (9)$$

Equation (9) is the asymmetry: with these inputs no optimizer on the allocation line can choose anything but $w_2 = 1 - w_1^{line} = 0.0340$. The residual FX openness is not neglect; it is the minimum-variance split of one notional envelope across an 18-to-1 variance mismatch with near-zero correlation. The budget then slides the European solution along the line to its vertex (upward in w_1 , because (4) makes w_1 the cheap direction and w_2 the expensive one): from 0.9660 to 0.9705. In the American regime there is no such slide. The cap is slack by KRW 8.93bn, so the optimum simply rests at the line minimum-variance point 0.9660 that (9) delivers, and the FX leg keeps the full 0.0340 of cover the variance arithmetic asks for. The two

regimes therefore share one unconstrained optimum and separate only through the budget.

6.3 Question (c): Cost Minimization Is a Different Program with a Different Answer

Because both cost functions are affine (4), pure cost minimization over the feasible polygon is a linear program whose optimum is a corner, identifiable by inspection: cost falls in w_1 and rises in w_2 , so the minimum is $(w_1, w_2) = (1, 0)$ (cover the oil leg completely, leave the currency leg completely open) in *both* regimes, confirmed by the solvers and by corner enumeration. The European cost minimum spends KRW 42.74bn (5.0% under budget) and accepts $\sigma_{res} = 0.0926$, which is +1.07% relative to the risk minimum; the American cost minimum spends KRW 33.43bn (25.7% under budget) at the same $\sigma_{res} = 0.0926$ (+1.08% relative), the two regimes coinciding in risk because at $(1, 0)$ the only surviving exposure is the FX leg, whose volatility is common to both.

The answer to question (c) is therefore an unambiguous *no*: “risk minimization subject to a budget” and “cost minimization” select different points in both regimes, a solution on the allocation line versus a polygon corner. The gap between them prices the last sliver of FX coverage: moving from the cost corner $(1, 0)$ to the risk solution trades KRW 2.26bn (European) or KRW 2.60bn (American) of additional spending for a 1.1% relative reduction in residual volatility: the entire remaining decision, compressed into one trade.

Re-anchoring the cost program with an explicit risk cap closes the gap: minimizing cost subject to $\sigma_{res} \leq \sigma_{res}^*$ (the achieved risk minimum) returns the risk-minimizing vertex itself in both regimes (Table 2). Economically this says the risk minima are *cost-efficient*: no cheaper allocation attains their risk level, so the budgeted risk program is already sitting on the cost—risk frontier. Methodologically it is the duality audit of Section 5 passing on live data.

6.4 The Global Geometry

Figure 2 assembles the whole problem on the full unit square: the residual-volatility surface (1) computed on a 201×201 grid, with all four constraints (6) evaluated at every node. Feasible faces are shaded (dark = low risk) and outlined in black; the constraint-violating remainder of the surface is drawn as a bare wireframe. Two properties are immediate. First, how little of the decision space survives the constraints: a sliver hugging $w_1 \gtrsim 0.965$ in the European regime, and a far roomier band rising from $w_1 = 0.843$ in the American regime, whose near-horizontal boundary slope of 0.0373 is the direct image of the KRW 4.22bn FX premium. Second, the risk minima (filled circles) sit where the shaded region touches its own darkest values: the minimum *within* the feasible shading, categorically distinct from the surface’s global minimum at the infeasible corner $(1, 1)$, where $\sigma_{res} = 0$ but the cost ledgers read KRW 55.1bn (European) and KRW 36.2bn (American) and the allocation envelope is violated outright. The cost-minimizing corner $(1, 0)$ (open triangles) is visibly a different point in both panels, and the zoomed bottom row resolves what the full-square view compresses: the shaded surface patch, its black boundary, and the two optima pinned to it by their drop stems. Figure 3 gives the same geometry in contour view with a zoom: the budget boundary (solid), the allocation line (dashed), the shaded feasible

Residual-volatility surface $\sigma_{res}(w_1, w_2)$: feasible region (all four constraints) shaded, dark = low risk; infeasible region wireframe only. Top: full unit square. Bottom: 3-D zoom.




Figure 2: Residual-volatility surface over the hedge-ratio square. Shaded faces (dark = low risk, black outline) satisfy all four constraints; the wireframe region violates at least one. Top row: full unit square. Bottom row: three-dimensional zoom onto the feasible neighborhood, where the vertex structure is legible: dotted stems drop from each optimum to the panel floor. Filled circle: risk minimum. Open triangle: cost minimum $(1, 0)$.

polygon, and the three solution markers per panel, making the vertex structure of Section 6.1 and the corner-versus-vertex contrast of Section 6.3 visible at a glance.

6.5 Parameter Robustness

Consider the worst case over an uncertainty box. Since σ_{res} is monotone increasing in σ_1 , σ_2 and ρ (all unhedged weights are nonnegative), the worst case is the upper corner of the box in closed form, and the min—max program replaces the objective with σ_{res} evaluated at inflated parameters. The feasible polygon is built entirely from prices, quantities and the budget, so it does not move; only the objective does. What happens next depends on how the box is drawn, and the distinction matters more than it first appears.

Iso-risk contours of σ_{res} , budget boundary (solid), $w_1 + w_2 = 1$ (dashed);
feasible region shaded gray. Top: full domain. Bottom: zoom.




Figure 3: Contour view of the same geometry. Iso- σ_{res} contours (gray), budget boundary (solid black), allocation line $w_1 + w_2 = 1$ (dashed), feasible region (shaded). Bottom row zooms onto the optima: the risk minimum (filled circle) sits at the vertex of the feasible polygon; the cost minimum (open triangle) at the corner $(1, 0)$; the risk-capped cost minimum ($\times$) coincides with the risk minimum: the duality audit passing.

Inflating the volatilities alone leaves both allocations untouched. Scaling σ_1 and σ_2 by a common factor leaves the variance ratio σ_1^2/σ_2^2 unchanged, and (9) depends on the two volatilities only through that ratio and through ρ . The line solution therefore stays at $w_1 = 0.9660$, which lies inside the European budget boundary at $w_1 = 0.9705$, so the European optimum remains the vertex $(0.9705, 0.0295)$ and the American optimum remains $(0.9660, 0.0340)$. Only the achieved risk re-prices, from 0.0916 to 0.1008. This is the sense in which a constraint-pinned optimum is robust: it is robust to the parameters that do not move the line.

Shifting the correlation moves both. Correlation is not such a parameter. Taking the box to be both volatilities inflated by 10% and ρ shifted up by 0.05, to $\rho = 0.1376$, the line solution (9) moves to

$$w_1^{line}|_{\rho=0.1376} = 0.9770, \quad (10)$$

which now lies *outside* the European budget boundary $w_1 \geq 0.9705$. The point $(0.9770, 0.0230)$ costs KRW 44.50bn, inside the KRW 45bn cap, so it is feasible and the cap goes slack. Both regimes therefore relocate to the same interior point $(0.9770, 0.0230)$, with achieved risk 0.1013 in each. The European vertex is not preserved under this box: the budget stops binding, and the optimum is pinned by curvature in both regimes rather than by the cap in one.

The contrast is the substantive result of this section, and it sharpens the limitation stated in Section 11. Vertex-pinning delivers robustness against proportional volatility mis-estimation, which is the error mode a desk usually has in mind, because such errors preserve the variance ratio that (9) actually reads. It delivers no robustness at all against correlation error, because ρ enters (9) separately and can push the line solution across the budget boundary, at which point the constraint that was doing the pinning simply stops binding. A shift of +0.05 in ρ , small next to the migration documented in stressed regimes, is already enough to do this. Estimation error therefore changes what the allocation *delivers* under every perturbation considered here, and changes what the allocation *is* whenever the perturbation moves ρ .

6.6 How much of the answer is the allocation constraint? A κ -sweep

Section 6.2 attributes the extreme w_1/w_2 split to the variance arithmetic. That attribution is conditional on $w_1 + w_2 \leq 1$, because (9) is derived on the line $w_1 + w_2 = 1$; it describes where the minimum sits *along* that line, not whether the line is where the optimum belongs. Relaxing the cap to $w_1 + w_2 \leq \kappa$ separates the two questions. Table 3 reports the risk-minimizing program at $\kappa \in [1, 2]$, holding budget, premiums and the stress ledger fixed.

Three things follow, and the third is a limitation rather than a result.

The FX leg is not intrinsically unattractive. At $\kappa = 1$ the European program buys 2.95% FX coverage. At $\kappa = 1.25$ it buys 18.31%, and residual volatility falls from 0.0916 to 0.0756, an improvement of 17.4% that no amount of extra budget can buy at $\kappa = 1$ (Section 6.1 showed the cap's shadow value exhausts at KRW 45.35bn). The 97/3 split is therefore not the statement that FX coverage is worth little. It is the statement that, when a barrel of coverage and a dollar of coverage compete for one notional envelope, the variance arithmetic prefers the barrel.

Under American premiums the constraint is the entire binding force. The two Shapley-attributed premiums sum to KRW 36.17bn at full coverage of both legs, comfortably inside the

κ	European				American			
	w_1	w_2	σ_{res}	cost	w_1	w_2	σ_{res}	cost
1.00	0.9705	0.0295	0.09160	45.00	0.9660	0.0340	0.09158	36.07
1.25	1.0000	0.1831	0.07562	45.00	0.9745	0.2755	0.06869	36.11
1.50	1.0000	0.1831	0.07562	45.00	0.9830	0.5170	0.04579	36.15
1.75	1.0000	0.1831	0.07562	45.00	0.9915	0.7585	0.02290	36.19
2.00	1.0000	0.1831	0.07562	45.00	1.0000	1.0000	0.00000	36.17

Table 3: Risk-minimizing allocation as the coverage cap $w_1 + w_2 \leq \kappa$ is relaxed. Costs in KRW bn against the KRW 45bn budget. At $\kappa = 1$ the reported optima of Section 6.1 are recovered. Relaxing κ changes the answer materially, and in the American regime removes the risk entirely.

KRW 45bn authority. At $\kappa = 2$ the American program therefore buys $w_1 = w_2 = 1$ and drives σ_{res} to zero, with the budget still slack by nearly KRW 9bn. Every interior result reported for the American regime in this paper, the 0.9660/0.0340 split included, exists only because $\kappa = 1$ forbids that corner.

The consequence for how the paper’s results should be read. The closed-form asymmetry (9) and the vertex geometry of Section 6.1 are correct statements about the program as specified, and the instrument-mix conclusion of Section 8, which compares two ways of buying the *same* WTI coverage, does not depend on κ at all. But the headline allocation is a joint consequence of the variance ratio and the coverage cap, and the cap carries at least as much of it as the arithmetic does. A firm whose treasury does not in fact regard the two legs as sharing one envelope, in the sense of (7), should not adopt the 0.97/0.03 split; it should re-solve at its own κ . This is a sharper caveat than the correlation sensitivity of Section 6.5, because it does not require the inputs to be mis-estimated at all.

6.7 A co-struck, co-termed respecification

The two regimes as specified price different strikes (5% in the money against at the money) and the two legs different maturities ($T_1 = 0.833$ against $T_2 = 0.5$), both inherited from the production configuration (Section 4). This subsection removes both mismatches and re-solves everything, as the last of the paper’s specification sweeps: both legs struck at the money and co-termed at $T = 0.833$, all other inputs unchanged.

Under the paper’s own pricing conventions the harmonized premiums are $P_{B76} = 10.91$ USD/bbl (full WTI cover KRW 35.58bn against 41.26bn) and $P_{GK} = 47.28$ KRW/USD (full FX cover KRW 7.90bn against 13.84bn), with marginal costs (-70.01 , $+6.42$) bn per unit ratio. Three consequences follow.

Question (a) reverses, by a fourth independent route. The harmonized European ledger costs KRW 37.06bn at (1,0) and KRW 39.66bn at the line optimum, so the KRW 45bn authority is slack in both regimes and the vertex exists only for $B \in [37.06, 39.66]$ bn. Both regimes sit at the line solution (0.9660, 0.0340): the allocation itself does not move, because (9) reads no contract parameter, and question (c)’s corner contrast survives unchanged in kind.

The instrument comparison flips sign, which strengthens the recommendation. The standalone knock-out leg of Section 8 is struck at the money; the vanilla leg it competes against is struck 5% in the money. Aligning the strikes, the same-strike vanilla costs KRW 35.58bn

against the knock-out's KRW 36.78bn: the knock-out is the *dearer* instrument, the indifference mortality is negative ($\bar{p}_{\text{harm}} = -1.1\%$), and the vanilla branch dominates at every mortality level. The current specification's positive switch point decomposes accordingly: its numerator $P_V - P_K = 4.48\text{bn}$ is a strike gap of 5.68bn, the intrinsic value the in-the-money vanilla charges for and the at-the-money knock-out does not carry, less a 1.20bn wedge by which the knock-out out-prices the same-strike vanilla. The $\bar{p} = 4.24\%$ of Section 8 therefore prices the strike mismatch, not a barrier discount.

What the wedge is, and is not. The 1.20bn is a cross-model quantity: the vanilla is a European Black-76 value at the raw volatility, the knock-out an American LSMC value under jump-diffusion at the split calibration, so American exercise value and the jump distribution sit on one side and the barrier discount on the other, and under these conventions the former outweigh the latter. The harmonization does not produce a clean market price of the barrier; what it establishes is the sign, that no positive knock-out discount survives strike alignment under this paper's pricing conventions, and hence that the all-vanilla selection of Section 8 holds *a fortiori*, at every mortality rather than above a threshold.

7 The Knock-Out Survival Haircut

This section and the next answer question (d). The American structure's premium advantage (KRW 31.95bn against KRW 41.26bn for full WTI cover) is bought with a knock-out barrier: the option dies if WTI trades outside its barriers before maturity. This creates a specific tension with the stress ledger of Section 3.2, which prices the *unhedged* fraction at the stress scenario while implicitly treating the *hedged* fraction as fully protected there. For a knock-out structure that treatment is optimistic in precisely the scenario that matters: at the stress price of 113 USD/bbl the upper barrier at 120 is close, and a knocked-out call protects nothing.

The exposure is priced with one parameter. Let p_{KO} be the probability that the structure is dead under the stress scenario. The protected fraction of each American leg is then $w(1 - p_{KO})$ and not w (one WTI barrier kills both legs of the joint structure), and the stress-adjusted total cost becomes

$$\tilde{C}^{AM}(w_1, w_2; p_{KO}) = \text{premiums} + (1 - w_1(1 - p_{KO})) UL_{WTI} + (1 - w_2(1 - p_{KO})) UL_{FX}, \quad (11)$$

where $UL_{WTI} = Q_{oil} \max(0, S_{WTI}^{st} - S_{WTI}) S_{KRW}^{st}$ and $UL_{FX} = Q_{USD} \max(0, S_{KRW}^{st} - S_{KRW})$ are the full-exposure stress losses, KRW 105.59bn and KRW 1.48bn respectively. At $p_{KO} = 0$, (11) reduces to (3).

Figure 4 traces (11) in p_{KO} , both at the adopted risk-minimizing allocation and minimized over all feasible (w_1, w_2) . Two structural facts emerge.

The budget is breached at $p_{KO}^ = 10.96\%$.* While $\partial \tilde{C}^{AM} / \partial w_1 < 0$ the cheapest attainable stress-adjusted cost is $\min_{(w_1, w_2)} \tilde{C}^{AM} = C^{AM}(1, 0) + p_{KO} UL_{WTI}$, which crosses the KRW 45bn cap at

$$p_{KO}^* = \frac{B - C^{AM}(1, 0)}{UL_{WTI}} = \frac{45,000,000,000 - 33,425,495,704}{105,586,000,000} = 0.1096. \quad (12)$$

Beyond a 10.9% stress knock-out probability, *no allocation whatsoever* keeps stress-adjusted cost within the budget: the term $p_{KO} \cdot UL_{WTI}$ is a loss floor that no choice of (w_1, w_2) can hedge away, because the instrument that would hedge it is the instrument that is dying.

That formula is valid only while the coefficient on w_1 in $\tilde{C}^{AM}$ is negative, that is while $Q_{oil}P_{WTI}^{Sh}(1 + r_w T_1) < (1 - p_{KO})UL_{WTI}$. The coefficient changes sign at

$$p_{KO}^\dagger = 1 - \frac{Q_{oil}P_{WTI}^{Sh}(1 + r_w T_1)}{UL_{WTI}} = 0.6974, \quad (13)$$

above which buying WTI coverage no longer pays for itself against the survival-adjusted stress loss and the unconstrained minimiser moves from $(1, 0)$ to $(0, 0)$, where the ledger reads $UL_{WTI} + UL_{FX} = \text{KRW } 107.06\text{bn}$. Since $p_{KO}^* = 0.1096$ lies below $p_{KO}^\dagger$, the infeasibility threshold is on the $(1, 0)$ branch and (12) applies. The held-contract mortality $p_{KO}^{\text{held}} = 0.4901$ established below also lies below $p_{KO}^\dagger$, so the cheapest attainable ledger there is the $(1, 0)$ branch value of KRW 85.2bn; only for mortalities above 0.6974 does the minimiser move to $(0, 0)$ and the floor flatten at $UL_{WTI} + UL_{FX} = \text{KRW } 107.06\text{bn}$. Figure 4 plots the branch-correct curve throughout.

Measuring p_{KO} : conditioning is everything. The structure's unconditional knock-out statistics already sit well above p_{KO}^* . Two are in circulation and they answer different questions: 45.1% is the raw barrier-touch frequency under exact first-passage monitoring, which is the correct unconditional statistic for a survival ledger, while 23.09% is that figure net of early exercise, which understates barrier risk because an exercised path is tallied as exercised rather than knocked out. The companion paper Park (2026b) makes the same point in its Section 6.1. The survival ledger of this section needs the barrier-touch figure, but they are the wrong numbers for the stress ledger: they average over paths launched from the calm spot of 78.94, most of which never approach the barrier. The number the ledger needs is conditional on the stress state itself. Re-simulating the calibrated jump-diffusion of Section 5 from $S_{WTI} = 113$ (only 6.2% below the 120 barrier, with 32.4% diffusive volatility, 6.8 jumps per year and 0.833 years still to run) gives

$$p_{KO}^{\text{stress}} = 0.8925 \pm 0.0007 \quad \begin{array}{l} \text{(200,000 paths, jump-adapted exact} \\ \text{first passage, asymmetric up/down} \\ \text{jumps),} \end{array} \quad (14)$$

2.0 times the unconditional barrier-touch rate and 8.2 times the infeasibility threshold (12).

Which conditional this is, and which one the ledger actually needs. Equation (14) launches a *fresh* contract from $S_{WTI} = 113$ with the full 0.833 years still to run and the barrier only 6.2% away, so it answers: if a desk were standing at the stress spot today and bought this structure, how likely is it to die? That is not the importer's question. The importer holds a contract struck and bought at 78.94 and wants to know whether that contract is still alive on the paths where WTI actually reaches the stress region. Measured that way on the same engine, conditioning the original contract on a terminal WTI in $[108, 118]$,

$$p_{KO}^{\text{held}} = 0.4901 \quad \text{(200,000 paths, 7,992 in the stress band),} \quad (15)$$

against an unconditional 0.4374 on the same bank. The held-contract figure is the one this

paper’s ledger should use, and the remainder of this section is read against it; (14) is retained because it is the quantity a desk pricing a *new* structure at the stress spot would want, and because the two together bracket the range over which the conclusion is tested.

The conclusion is unchanged by the substitution, and that is the point of reporting both. At $p_{KO}^{\text{held}} = 0.4901$ the all-KO ledger reads KRW 86.04bn against the KRW 45bn authority, still nearly twice the cap; the mortality is 11.6 times the instrument-switch indifference point $\bar{p} = 4.24\%$ of Section 8 and 4.5 times the infeasibility threshold $p_{KO}^* = 10.96\%$. Every threshold in Sections 7 and 8 is crossed on either measure. What changes is the headline magnitude: the honest statement is that an all-KO book is dead on roughly half the paths that reach the stress region, not on nine-tenths of them. The jump calibration underlying this simulation is itself asymmetric, fitted separately as an up-jump regime ($\lambda_{up} = 2.328/\text{yr}$, $\theta_{up} = +8.04\%$, $\delta_{up} = 1.68\%$) and a down-jump regime ($\lambda_{dn} = 4.462/\text{yr}$, $\theta_{dn} = -8.75\%$, $\delta_{dn} = 2.78\%$) from the underlying return history, and not the single symmetrized pool ($\lambda, \theta_J, \delta_J$) that the pricing engines elsewhere in this paper consume for tractability. Re-deriving (14) under that symmetrized pool instead gives 0.8937 ± 0.0007 , a tenth of a percentage point away, because the 32.4% diffusive volatility dominates barrier-touch dynamics over a sub-annual horizon and the jump distribution’s directional asymmetry is a second-order contributor to this specific question. On either measure the pure-KO book’s arithmetic is unambiguous (Figure 4): the adopted allocation’s ledger reads KRW 86.04bn at p_{KO}^{held} and KRW 127.10bn at p_{KO}^{stress} , against a KRW 45bn authority in both cases. Under the very scenario the hedge is sized against, an all-KO book is, on roughly half the paths that get there, no hedge at all.

For the pure-KO book the gap is a reporting item; Section 8 makes it a decision. The allocation program of Section 3 deliberately keeps (3), not (11), in its budget constraint: constraining on (11) would render the single-instrument program infeasible without changing any decision margin, since the floor $p_{KO} UL_{WTI}$ is allocation-invariant. What *can* respond to (11) is the choice of instrument, which is the degree of freedom the next section adds.

8 The Mixed Vanilla/Knock-Out Program

The haircut analysis treats the instrument as given and finds that no coverage choice can rescue an all-KO book under stress. The natural response is to stop treating it as given. This section lets the optimizer choose the instrument along with the coverage: the WTI book is split as $w_1 = w_1^V + w_1^K$ between a vanilla Black-76 leg (which survives stress) and an independently-priced KO leg (which dies with probability p_{KO}), the FX leg is priced with the vanilla Garman—Kohlhagen premium, and, decisively, the *survival-adjusted* ledger becomes the budget constraint. The KO leg’s premium P_K here is *not* the Shapley-attributed WTI share (KRW 15,093.75/bbl) used to cost the American structure everywhere else in this paper: that figure is a cooperative-game marginal-contribution split of the *jointly*-priced quanto premium, valid as a cost-reporting convention for a single indivisible structure but not the market price of a standalone, independently-tradable WTI KO call. Once the WTI book is genuinely split into two separately-tradable legs, the KO leg needs its own price: an independent single-asset LSMC American KO call, priced with no FX coupling and no cost attribution on 200,000 jump-adapted




Figure 4: Stress-adjusted total cost as a function of the stress knock-out probability p_{KO} : the pure-KO book at the adopted allocation (solid) and at its cheapest attainable allocation (dashed), against the mixed vanilla/KO program of Section 8 (dash-dotted), which is ledger-feasible for every p_{KO} . The dashed curve switches branch at $p_{KO}^\dagger = 0.6974$ (13), above which the unconstrained minimiser moves from $(1, 0)$ to $(0, 0)$. Small markers: unconditional knock-out statistics. Star: the relaunched-contract measurement (14). The held-contract mortality (15) of 0.4901, which Section 7 argues is the figure the ledger needs, sits at KRW 86.1bn.

exact-first-passage paths, comes to $P_K = \text{KRW } 17,377.11/\text{bbl}$, 15.1% above the Shapley share. Both premiums in this section are cash paid at inception and both are carried to the decision horizon on the same simple-interest factor $(1 + r_w T_1)$ used throughout. Every number in this section uses that standalone price:

$$\min_{w_1^V, w_1^K, w_2} \sigma_{res}(w_1^V + w_1^K, w_2) \quad \text{s.t.} \quad \tilde{C}^{mix} \leq B, \quad w_1^V + w_1^K + w_2 \leq 1, \quad w \geq 0, \quad (16)$$

$$\tilde{C}^{mix} = P_V(w_1^V) + P_K(w_1^K) + P_{GK}(w_2) + (1 - w_1^V - w_1^K(1 - p_{KO}))UL_{WTI} + (1 - w_2)UL_{FX}. \quad (17)$$

Three structural facts make (16) tractable and its answer sharp.

First, the program is feasible for every p_{KO} . The all-vanilla corner $(w_1^V, w_1^K, w_2) = (1, 0, 0)$ costs KRW 42.74bn regardless of p_{KO} (vanilla protection does not die), so the infeasibility that kills the all-KO book (12) simply cannot occur. The off-ledger gap of Section 7 becomes an on-ledger constraint the optimizer can actually satisfy.

Second, the split is bang-bang, with a closed-form indifference point. At any fixed total coverage $w_1^V + w_1^K$, (17) is linear (affine) in the split, so its two marginal costs are constants, read directly off (17) by differentiating the KO- and vanilla-coverage terms separately:

$$\frac{\partial \tilde{C}^{mix}}{\partial w_1^K} = P_K - (1 - p_{KO})UL_{WTI}, \quad \frac{\partial \tilde{C}^{mix}}{\partial w_1^V} = P_V - UL_{WTI}. \quad (18)$$

The first says: buying one more unit of KO coverage costs its premium P_K but only removes the stress loss with probability $1 - p_{KO}$, so the marginal saving is discounted by survival odds.

The second says: vanilla coverage costs P_V but removes the stress loss with certainty. Since a linear program's optimal split is always a corner and not an interior mix, the two instruments are indifferent where their marginal costs cross, $\partial \tilde{C}^{mix} / \partial w_1^K = \partial \tilde{C}^{mix} / \partial w_1^V$:

$$P_K - (1 - \bar{p}) UL_{WTI} = P_V - UL_{WTI} \implies \bar{p} = \frac{P_V - P_K}{UL_{WTI}}, \quad (19)$$

$$\bar{p} = \frac{41,258,998,746 - 36,780,738,568}{105,586,000,000} = 0.0424. \quad (20)$$

Below $\bar{p}$, (18)'s first expression is the smaller (more negative) of the two (KO coverage is the cheaper marginal unit, the discount outrunning its mortality), and the optimizer fills w_1^K first, so the book is all-KO; above $\bar{p}$ the ranking flips, the discount no longer pays for its own mortality, and it fills w_1^V first, so the book is all-vanilla. The optimal split is therefore an endpoint, switching at $\bar{p}$, as a linear program's corner-solution property requires; at $\bar{p}$ itself the two branch cost functions coincide at every allocation, so the switch in Figure 5 occurs at a single point. Because the standalone KO price sits well above the Shapley-attributed joint-quanto share, the switch point is low: a KO leg priced at its true independent cost is worth comparatively little mortality risk before the vanilla alternative dominates. The level of $\bar{p}$ also carries a specification component: $P_V - P_K$ compares a 5%-in-the-money vanilla with an at-the-money knock-out, and Section 6.7 shows that aligning the strikes turns the difference negative, so the positive switch point prices the strike gap rather than a barrier discount, and the all-vanilla selection below only strengthens under alignment.

Third, the solution has three regimes (Figure 5). For $p_{KO} \leq 0.0390$ the all-KO branch is cheap enough on the survival-adjusted ledger that the budget goes slack and the program reaches the unconstrained variance floor (0.9660, 0.0340), $\sigma_{res} = 0.0916$, the same floor the unadjusted American program of Section 6.1 attains under a KRW 45bn authority. For $0.0390 < p_{KO} < \bar{p}$ the all-KO branch is budget-pinned and its achieved risk deteriorates as mortality eats the ledger. At $\bar{p} = 0.0424$ the book switches to all-vanilla and stays there: allocation (0.9705, 0.0295), survival-adjusted cost KRW 45bn, $\sigma_{res} = 0.0916$, invariant in p_{KO} . The pure-KO death threshold, recomputed with the same standalone KO price, is $p^* = 0.0638$, well above the switch: the optimizer abandons the KO leg at $\bar{p} = 0.0424$, while an all-KO book remains budget-feasible up to 0.0638. The instrument is dropped on its economics well before it becomes unaffordable, and the interval between the two thresholds is the range in which a firm could still buy the knock-out but should not.

At the measured stress mortality $p_{KO}^{stress} = 0.8925$ the verdict is not close: the mixed program selects the all-vanilla book, roughly twenty-one times $\bar{p}$ deep into the vanilla regime. Two comparisons calibrate what is gained. Against the pure-KO book at the same p_{KO} , the mixed optimum's ledger is KRW 45.00bn versus 130.67bn on the standalone-priced ledger: the entire KO-mortality exposure, KRW 86bn, priced out by the instrument switch. And against the pure-KO program's own optimum on the *unadjusted* ledger, the mixed optimum matches its residual volatility, 0.0916, while satisfying a strictly harder constraint. The instrument switch therefore costs no variance: survival-adjusted feasibility comes free. That comparison is in fact conservative, because the risk objective (1) is instrument-blind: it credits w_1 with full protection

whether the WTI leg is a vanilla call or a knock-out that is dead with probability p_{KO} . Sections 7 and 8 correct for mortality on the cost ledger alone and leave the objective untouched. Carrying the same correction into the objective, by crediting the knock-out leg with effective coverage $w_1(1 - p_{KO})$, raises the all-KO book’s residual volatility at the measured mortality from 0.0916 to 0.372, against 0.0916 for the vanilla book. Under a survival-adjusted objective the vanilla structure is not merely equal-risk at lower survival-adjusted cost but strictly dominant on both axes; the ledger-only treatment reported here understates the case for the switch. The knock-out structure’s KRW 4.48bn premium discount over the standalone vanilla leg, the only argument in its favor, is the quantity (19) says is overwhelmed once stress mortality exceeds 4.2%; the measured value is 49%, a factor of twelve beyond the point at which the discount stops paying for its own mortality.




Figure 5: The mixed program’s optimal residual volatility as a function of p_{KO} , with its two branches, using the independently-priced standalone KO premium. Floor regime ($p_{KO} \leq 0.0390$): the all-KO ledger leaves the budget slack and the unconstrained variance floor $\sigma_{res} = 0.091585$ is attained. Budget-pinned KO regime (0.0390—0.0424): achieved risk deteriorates as mortality consumes the ledger. All-vanilla regime ($p_{KO} \geq \bar{p} = 0.0424$): the book switches instruments and the optimum is invariant at $\sigma_{res} = 0.091602$. The held-contract stress mortality $p_{KO}^{\text{held}} = 0.4901$ lies far inside the vanilla regime. The vertical axis is scaled to the 1.8×10^{-5} band the whole switch occupies, so the step at $\bar{p}$ is legible; on a conventional axis the two levels are indistinguishable, which is itself the point that the instrument switch costs essentially no variance. The all-KO branch leaves the frame immediately above $\bar{p}$.

9 Discussion

Five points deserve emphasis beyond the numerical answers.

The budget’s main effect is on governance and not on risk. Section 6.1 shows the cap costs almost no residual volatility where it binds at all, yet it converts an open-ended optimization into a vertex with an auditable location: the point where the combined ledger exhausts authority and coverage exhausts the notional envelope. For a treasury committee, “we are at the corner of what you authorized” is a materially different (and more defensible) statement than “we

are at an interior optimum of a model you have not seen.” The shadow prices computed here give that committee the exchange rate between authority and risk: roughly half a basis point of residual volatility per additional billion KRW in the European regime, decaying to zero at KRW 45.35bn, and zero throughout in the American, where the optimal programme already fits inside the cap at KRW 36.02bn.

The asymmetry follows from the risk objective and not from the budget constraint. It is tempting to read $w_2 \approx 0.03$ as the FX hedge being “crowded out” by the budget. Equation (9) suggests a different reading: even with the budget deleted, the minimum-variance split leaves the FX leg 96.6% open. What sets the split is the variance ratio $\sigma_1^2/\sigma_2^2 \approx 18.2$ together with the low estimated correlation, both properties of the sample used for calibration and not policy choices. An importer who wants more FX coverage must argue for a different objective (for example, a tail metric on the KRW cash outflow) and not for a bigger budget.

The closed-form split inherits whatever the correlation estimate assumes. The 0.97/0.03 allocation is a consequence of (9) evaluated at a single unconditional correlation, $\rho = 0.0876$, estimated over 1,299 daily observations of a predominantly calm sample. Because the two legs are close to orthogonal at that value, the variance objective sees almost no diversification benefit from the FX leg and assigns it almost nothing. This dependence is not incidental. In stressed regimes, cross-asset correlations are widely documented to migrate sharply toward ± 1 as common liquidity and risk-appetite factors dominate idiosyncratic ones, a phenomenon usually described as correlation breakdown. Under such a migration the term $2w_1w_2\rho\sigma_1\sigma_2$ in (1) stops being negligible, the minimum-variance line rotates, and the interior split that (9) returns moves away from the corner reported here. The static single- ρ formulation is therefore best read as a calibration to normal-market co-movement; it is precisely in the joint tail, where an oil-price spike and a KRW depreciation are most likely to arrive together, that the allocation reported here has the least claim to optimality. Section 11 states the consequence.

Premium regimes should be compared on the survival-adjusted ledger, which now has a decision attached. On the naive ledger the American structure dominates: full WTI cover for KRW 31.95bn against KRW 41.26bn, an apparent saving of KRW 9.3bn. That comparison, however, prices the American leg at its Shapley-attributed share of the jointly-priced quanto structure; the independent, standalone-tradable KO instrument that the mixed program of Section 8 actually offers against a standalone vanilla leg costs KRW 36.78bn against the same KRW 41.3bn, a saving of only KRW 4.48bn. Sections 7 and 8 show that saving, whichever way it is measured, is bought with mortality, and Section 6.7 shows it is not even a barrier discount, vanishing outright once the two instruments are struck alike; the mortality price is currently ruinous: conditioning the knock-out probability on the stress state lifts it from the unconditional 23% to a measured 49% for the contract actually held, at which the all-KO ledger reads KRW 86.04bn against a 45bn authority. The European vanilla hedge is expensive because it survives; the American knock-out hedge is cheap because it very probably will not. The mixed program turns this from commentary into allocation: it prices the trade at a single indifference number, $\bar{p} = 4.24\%$ of stress mortality, and at roughly twenty-one times that number it walks the book to vanilla on its own.

The vanilla recommendation of Section 8 operates at a different layer than the dynamic-

hedging analysis of this same structure in Park (2026b), and the two are not in tension. That analysis holds the actual barrier contract and asks how its hedge ratio should be computed. Its controlled findings are that the apparently large cost of a barrier-blind proxy delta is dominated by payoff, premium and stopping-time differences between the engines being compared rather than by delta quality; that what remains on a like-for-like ledger is reproduced by a constant hedge ratio carrying no barrier information; and that the currency leg should be sized to the structure's own carrying value V/S_2 , not to a multiple of the WTI delta. The instrument-mix program of Section 8 is a different decision, made earlier: it chooses *which contract to buy in the first place*, before any hedging begins. When the measured stress mortality selects the vanilla branch, the position that results is a barrier-free WTI call correctly hedged by the matching vanilla delta, and the model-instrument mismatch the companion paper dissects never arises. A firm that took both papers' conclusions simultaneously would not be computing barrier deltas at all: it would, per Section 8, have already exited the barrier position in favor of the instrument whose delta model is exact.

10 Practical Implications

For the importer's hedging desk the results reduce to four actionable statements, of which the first is the operative one. (1) Under the measured stress knock-out probability of 49% for the contract actually held, the WTI book should be held in the vanilla instrument: the KO structure's premium discount is only worth taking if stress mortality is believed below $\bar{p} = 4.24\%$, more than an order of magnitude under the measurement, the discount disappears entirely once the instruments are struck alike (Section 6.7), and the mixed program of Section 8 (feasible at every mortality level, budget-binding at KRW 45bn, and matching the pure-KO optimum's residual volatility) is the implementable form of the recommendation. (2) The adopted allocations are optima of a convex program, obtained from two independent solvers and cross-checked by re-solving the cost program under a risk cap; in the European regime a proportional 10% volatility mis-estimate leaves the allocation itself unchanged, because it preserves the variance ratio (9) reads, while a correlation shift of +0.05 relocates both regimes to (0.9770, 0.0230) with the cap slack (Section 6.5); in the American regime, pinned by the allocation line and not the cap, the same mis-estimate shifts it only from 0.9660 to 0.9770 (Section 6.5). (3) In the European regime as specified the assumed KRW 45bn ledger is fully committed, of which KRW 40.45bn is premium and KRW 4.55bn is the booked stress loss, though a ledger larger by 0.8% would not bind at all (8), and raising it to KRW 45.35bn captures the entire remaining variance benefit; in the American regime the optimal programme spends KRW 36.02bn, so the cap is not the operative constraint and further authority buys nothing. (4) The FX leg's near-zero coverage is optimal under the variance objective; if that outcome is commercially unacceptable, what must change is the objective and not the optimizer.

11 Limitations

The risk objective is a one-period standard deviation built from historical volatilities and a single correlation estimate, and it is instrument-blind: it credits w_1 with full protection whether the WTI leg is a vanilla call or a knock-out, which is why the American rows of Table 2 report the program’s objective value rather than the risk the American book carries (Section 8). Premium constants are treated as invariant to (w_1, w_2) . The static problem fixes ratios once; subsequent dynamic management is the subject of the companion paper. Three boundaries are quantitative and are stated with their size.

Correlation is the binding uncertainty. Section 6.5 shows that proportional volatility inflation preserves the variance ratio and hence the split, while a shift in ρ does not, and Section 6.5’s standard error puts the 95% interval on $\hat{\rho} = 0.0876$ at $[0.034, 0.142]$, mapping to FX coverage of $[2.21\%, 4.54\%]$ against a reported 3.40%. The vertex that answers question (a) requires $\rho < 0.1083$, three quarters of a standard error above the estimate. In stressed regimes correlations migrate toward ± 1 , a pattern documented for equity markets by Longin and Solnik (2001) and Ang and Chen (2002) and reported across asset classes since, and a WTI spike with a sharp KRW depreciation is precisely the joint event this hedge is bought against. A regime-switching covariance, a worst-case- ρ robust program in the sense of Ben-Tal and Nemirovski (1998) and Goldfarb and Iyengar (2003), or a tail objective such as the conditional value-at-risk of Rockafellar and Uryasev (2000), coherent in the sense of Artzner et al. (1999), would each be a natural extension and each would be expected to shift weight toward the FX leg.

The American premium split is a convention. Section 3.3: the companion paper’s factorization makes $v(\{1\}) = v(\{1, 2\})$, so the Shapley FX share is half a deterministic intrinsic value. Totals and the budget-slack answer are invariant to the attribution; marginal FX-cost statements and the Section 7 threshold (12) are not, and the decisions of Section 8 are made on standalone prices for exactly that reason.

The coverage cap carries much of the answer. Section 6.6 sweeps $w_1 + w_2 \leq \kappa$ and shows the reported split is a joint consequence of the variance ratio and the cap. Under American premiums at $\kappa = 2$ the program buys full cover of both legs for KRW 36.17bn with $\sigma_{res} = 0$. The cap’s rationale is the notional identity (7), which holds at the calibration spot; a firm that treats the two legs as separate books should re-solve at its own κ .

Question (a) is conditional; questions (b), (c) and (d) are not. Four independent conditions each reverse the budget-geometry result of Section 6.1: a coverage cap $\kappa > 1$ (Section 6.6), an authority above KRW 45.35bn (8), a market-convention contract specification, under which full European WTI cover costs KRW 4.86bn rather than 41.26bn and the cap is slack in both regimes (Section 4), and the co-struck, co-termed respecification of Section 6.7. None of the three touches the closed-form asymmetry, which reads only the variance ratio and ρ ; the cost-versus-risk contrast, which compares corners of one ledger; or the instrument-mix switch, whose indifference point is a ratio of premium differences and keeps its sign under re-specification. The paper’s first question is the fragile one and its later ones are not.

A monthly budget against a ten-month contract. B is described as a monthly authority while $T_1 = 0.833$ years, so in steady state ten tranches overlap and the KRW 45bn is ambiguous between spend per month and spend per position. The static one-shot program studied here does

not have to resolve that, because it allocates a single tranche once, but a rolling implementation would, and the overlapping-tranche version of this problem is not solved in this paper.

The spending authority is assumed. The KRW 45bn budget has no external source; Section 4 states it as 18.5% of the monthly crude bill. The European vertex requires $B \in [42.74, 45.35]$ bn by (8), and the assumed value sits 0.77% inside the upper endpoint, so question (a)’s answer is the least robust result in the paper: a budget larger by less than one percent reverses it. The American regime is unaffected, since its cap is slack at KRW 36.02bn throughout the interval and well beyond.

One stress point. The cost functions price the unhedged exposure at (113, 1550) rather than over a distribution. The vertex of Section 6.1 requires a stress WTI near that level and is gone by 95, where the European ledger costs KRW 43.45bn at the line optimum and the cap goes slack; at a stress USD/KRW of 1700 no allocation is feasible in either regime. Both thresholds scale inversely with the assumed stress loss, $\bar{p}$ moving from 23.8% at a stress WTI of 85 to 2.0% at 150. The direction of the instrument-mix switch does not depend on the stress point; its magnitude does.

The closed-form asymmetry and the vertex geometry are structural properties of the convex program. The quantitative allocation is conditional on the covariance inputs and on κ , and should be read against the intervals above rather than at the four decimal places in which it is reported.

12 Conclusion

A fixed-budget importer choosing static WTI and FX hedge ratios faces a problem whose solution geometry can be mapped in closed form and verified numerically, as this paper has done; its central finding is about the instrument, and the geometry is the machinery that prices it.

The American structure’s premium advantage carries a mortality that this paper measures rather than assumes. Re-simulated from the stress state itself with jump-adapted exact first-passage monitoring on 200,000 paths under the genuine asymmetric up/down jump calibration (the symmetrized-pool re-derivation lands a tenth of a percentage point away, so the simplification the pricing engines rely on is not driving the result), the knock-out probability for a contract newly struck at the stress spot is 0.8925, and for the contract actually held it is 0.4901 conditional on WTI reaching the stress region, the latter still 4.5 times the level at which an all-KO book’s survival-adjusted ledger can fit any allocation inside the budget. The cheap hedge’s protection is absent on roughly half the paths that reach the scenario it is bought against. The mixed vanilla/knock-out program turns this from commentary into allocation: priced with an independently-derived standalone KO premium rather than the jointly-priced structure’s Shapley attribution, feasible at every mortality level, switching instruments bang-bang at a closed-form indifference point of 4.24%, it walks the book to all-vanilla at the measured mortality and matches the pure knock-out optimum’s residual volatility under a strictly harder constraint. And the switch point itself is generous to the knock-out: struck like for like, the knock-out costs more than the vanilla under this paper’s pricing conventions (Section 6.7), so

the instrument’s one argument, its discount, prices the strike gap rather than the barrier, and the vanilla book dominates at every mortality.

Around that result, the allocation geometry. Risk minimization drives coverage to the allocation envelope in both regimes, and to the budget boundary in one: $(w_1^*, w_2^*) = (0.9705, 0.0295)$ under European vanilla premiums, committing the full assumed KRW 45bn ledger at a vertex that exists only over a 6.1%-wide band of budgets (8) and is undone by any of four independent respecifications, and $(0.9660, 0.0340)$ under American knock-out premiums, spending KRW 36.02bn with the cap slack, each identified as the global optimum of a convex program by two independent solvers and a dual check. Where the budget binds it does so shallowly, with a shadow price of about half a basis point of residual volatility per billion KRW and a hard floor at KRW 45.35bn. The allocation asymmetry, 0.97/0.03 in both regimes, is a closed-form consequence of an 18-to-1 variance ratio and a 0.09 correlation and not a preference of the optimizer, and it inherits that correlation estimate’s calm-sample provenance: under the correlation breakdown characteristic of stressed regimes the same closed form would return a different split. Cost minimization is a genuinely different program that runs to the corner $(1, 0)$ and underspends the budget by up to a quarter, converging with the risk program only when re-anchored by the risk cap, at which point the coincidence is arithmetic rather than evidential, since the re-anchored cost program has a unique feasible point. The static allocation question that precedes every dynamic hedging choice turns out, for this importer, to have an answer that is asymmetric, instrument-aware, and conditional exactly where the paper says it is: on the covariance structure, the coverage cap, and the contract specification on which the budget geometry is built.

Code and data availability. Every engine, the frozen derived inputs, and the scripts that re-generate each table and figure in this paper are available at github.com/JihongParker/wti-fx-hedge-program

References

- [1] Adler, M., and B. Dumas (1984). “Exposure to Currency Risk: Definition and Measurement.” *Financial Management* 13(2), 41—50.
- [2] Ahn, D.-H., J. Boudoukh, M. Richardson, and R. F. Whitelaw (1999). “Optimal Risk Management Using Options.” *Journal of Finance* 54(1), 359–375.
- [3] Anderson, R. W., and J.-P. Danthine (1981). “Cross Hedging.” *Journal of Political Economy* 89(6), 1182—1196.
- [4] Ang, A., and J. Chen (2002). “Asymmetric Correlations of Equity Portfolios.” *Journal of Financial Economics* 63(3), 443–494.
- [5] Artzner, P., F. Delbaen, J.-M. Eber, and D. Heath (1999). “Coherent Measures of Risk.” *Mathematical Finance* 9(3), 203–228.
- [6] Ben-Tal, A., and A. Nemirovski (1998). “Robust Convex Optimization.” *Mathematics of Operations Research* 23(4), 769–805.

- [7] Black, F. (1976). “The Pricing of Commodity Contracts.” *Journal of Financial Economics* 3(1—2), 167—179.
- [8] Brown, G. W., and K. B. Toft (2002). “How Firms Should Hedge.” *Review of Financial Studies* 15(4), 1283–1324.
- [9] Campello, M., C. Lin, Y. Ma, and H. Zou (2011). “The Real and Financial Implications of Corporate Hedging.” *Journal of Finance* 66(5), 1615–1647.
- [10] Carr, P., and A. Chou (1997). “Breaking Barriers.” *Risk* 10(9), 139–145.
- [11] Chopra, V. K., and W. T. Ziemba (1993). “The Effect of Errors in Means, Variances, and Covariances on Optimal Portfolio Choice.” *Journal of Portfolio Management* 19(2), 6–11.
- [12] DeMiguel, V., L. Garlappi, and R. Uppal (2009). “Optimal Versus Naive Diversification: How Inefficient Is the 1/N Portfolio Strategy?” *Review of Financial Studies* 22(5), 1915–1953.
- [13] Derman, E., D. Ergener, and I. Kani (1995). “Static Options Replication.” *Journal of Derivatives* 2(4), 78–95.
- [14] Ederington, L. H. (1979). “The Hedging Performance of the New Futures Markets.” *Journal of Finance* 34(1), 157—170.
- [15] Froot, K. A., D. S. Scharfstein, and J. C. Stein (1993). “Risk Management: Coordinating Corporate Investment and Financing Policies.” *Journal of Finance* 48(5), 1629—1658.
- [16] Garman, M. B., and S. W. Kohlhagen (1983). “Foreign Currency Option Values.” *Journal of International Money and Finance* 2(3), 231—237.
- [17] Gilje, E. P., and J. P. Taillard (2017). “Does Hedging Affect Firm Value? Evidence from a Natural Experiment.” *Review of Financial Studies* 30(12), 4083–4132.
- [18] Goldfarb, D., and G. Iyengar (2003). “Robust Portfolio Selection Problems.” *Mathematics of Operations Research* 28(1), 1–38.
- [19] Jin, Y., and P. Jorion (2006). “Firm Value and Hedging: Evidence from U.S. Oil and Gas Producers.” *Journal of Finance* 61(2), 893—919.
- [20] Longin, F., and B. Solnik (2001). “Extreme Correlation of International Equity Markets.” *Journal of Finance* 56(2), 649–676.
- [21] Longstaff, F. A., and E. S. Schwartz (2001). “Valuing American Options by Simulation: A Simple Least-Squares Approach.” *Review of Financial Studies* 14(1), 113—147.
- [22] Maillard, S., T. Roncalli, and J. Teiletche (2010). “The Properties of Equally Weighted Risk Contribution Portfolios.” *Journal of Portfolio Management* 36(4), 60–70.
- [23] Markowitz, H. (1952). “Portfolio Selection.” *Journal of Finance* 7(1), 77—91.

- [24] Merton, R. C. (1976). “Option Pricing When Underlying Stock Returns Are Discontinuous.” *Journal of Financial Economics* 3(1—2), 125—144.
- [25] Michaud, R. O. (1989). “The Markowitz Optimization Enigma: Is Optimized Optimal?” *Financial Analysts Journal* 45(1), 31–42.
- [26] Park, J. (2026b). “The Currency Leg of a Knock-Out Oil Option under Jump-Diffusion: Factorization, Netting, and What a Fitted LSMC Delta Does Not Buy.” Working paper, Pusan National University.
- [27] Pérez-González, F., and H. Yun (2013). “Risk Management and Firm Value: Evidence from Weather Derivatives.” *Journal of Finance* 68(5), 2143–2176.
- [28] Rampini, A. A., A. Sufi, and S. Viswanathan (2014). “Dynamic Risk Management.” *Journal of Financial Economics* 111(2), 271–296.
- [29] Reiner, E. (1992). “Quanto Mechanics.” *Risk* 5(3), 59—63.
- [30] Rockafellar, R. T., and S. Uryasev (2000). “Optimization of Conditional Value-at-Risk.” *Journal of Risk* 2(3), 21–41.
- [31] Roncalli, T. (2013). *Introduction to Risk Parity and Budgeting*. Chapman & Hall/CRC.
- [32] Shapley, L. S. (1953). “A Value for n -Person Games.” In *Contributions to the Theory of Games II*, Princeton University Press, 307—317.
- [33] Stulz, R. M. (1996). “Rethinking Risk Management.” *Journal of Applied Corporate Finance* 9(3), 8—25.