

IFRS 9 Hedge Accounting for a Quanto Knock-Out Structure

Jihong Park

Pusan National University

bjihong935@gmail.com

July 2026

Abstract

A Korean crude-oil importer's commodity-and-currency exposure can be hedged under IFRS 9 either as one combined instrument (a quanto knock-out call) or as two independent risk-component designations. This paper builds a full path-wise cash-flow-hedge accounting engine to quantify how the two architectures differ on hedge ineffectiveness, OCI dynamics, and the economic risk surviving a knock-out event. The combined designation's mean ineffectiveness comes out roughly three-and-a-half times larger than the split designation's, almost entirely from the split structure's commodity leg and not its currency leg, while full-horizon economic risk differs by only 1.3% of level. Extending the comparison to barrier-free variants of both designations isolates the driver: the split leg's ineffectiveness falls to a level indistinguishable from zero regardless of the barrier, while the combined structure's stays several times larger in the barriered configuration (24.78 against 7.70bn, a factor of 3.2) and does not approach zero in the barrier-free one, so what drives its elevated ineffectiveness is multiplying two floating exposures into one payoff, not the knock-out feature itself. Every headline figure is traced directly to the accounting engine's own live computation, and the SFP balance is confirmed at every step, though by construction of the pro-forma rather than as an independent test; the flow-level decomposition that would be an independent test instead fails, by an amount equal to the reclassification the knock-out sweep produces, and is reported as a result. The paper argues that the ineffectiveness gap is better understood as a consequence of IFRS 9's hypothetical-derivative convention than as a difference in economic risk.

Keywords: IFRS 9; cash flow hedge; hedge effectiveness; hypothetical derivative; designation architecture; knock-out barrier; FVTPL reclassification

JEL classification: M41; G32; G13; M48

Contents

1	Introduction	4
2	Literature Review	5

3	Product and Pricing Definition	6
3.1	Correlated jump-diffusion dynamics	6
3.2	Double knock-out barrier	7
3.3	Quanto payoff construction and the correlation coupling	8
3.4	Longstaff—Schwartz pricing surface and backward-induction algorithm	8
4	Designation Architectures	9
4.1	Structure A: single quanto KO designation	9
4.2	Structure B: dual independent designations	9
5	The Hypothetical Derivative	9
6	Accounting Engine	10
6.1	Signed lower-of test	10
6.2	Cost-of-hedging amortization	10
6.3	Knock-out event processing	10
6.4	Dynamic hedge ratio and non-circularity	11
6.5	Numerical safeguards	11
6.6	Per-step processing sequence	12
6.7	Maturity treatment	12
7	Economic Layer	12
7.1	Step residuals and telescoping	12
7.2	Post-KO naked exposure	13
8	Empirical Results	13
8.1	Inception and pricing-surface validation	13
8.2	H1: cumulative ineffectiveness	14
8.3	H2: full-horizon economic residual	14
8.4	H3: post-KO naked exposure and reclassification	15
8.5	H4: SFP and disclosure footprint	16
8.6	The maturity-freeze mechanism	18
9	Sensitivity	20
9.1	A validated fourth cell: the combination drives the gap	21
10	Discussion	23
10.1	What the floor and the clip do to the reported numbers	25
11	Limitations	26
12	Conclusion	27
A	Mathematical Derivations	27
A.1	Derivation of the hypothetical-derivative closed form	27

1 Introduction

A Korean crude-oil importer with a fixed monthly physical requirement of barrels= 2,000,000 (equivalently a USD funding need of roughly \$157.88 million per month at inception) is exposed to two risk factors that enter the KRW cost of imported crude *multiplicatively*: the US-dollar price of WTI, S_1 , and the USD/KRW exchange rate, S_2 . Neither risk factor alone determines the cash outflow; the product $S_1 \cdot S_2$ does. IFRS 9 explicitly recognizes this kind of joint exposure under the heading of an *aggregated exposure* (IFRS 9.6.3.4): a combination of a commodity price risk component and a foreign-currency risk component that together create a single, differently-shaped exposure, which may itself be designated as the hedged item in a cash-flow hedge.

That single aggregated-exposure framing admits two structurally different accounting architectures, and this paper’s entire empirical exercise is to measure, not assume, how they differ:

- **Structure A** designates a single hedging instrument, a double knock-out (KO) quanto call whose payoff already multiplies the WTI intrinsic value by the floating terminal FX rate, against the aggregated exposure as a whole. One hedging instrument, one hedge relationship.
- **Structure B** instead exercises IFRS 9’s separate permission (9.6.3.1) to designate hedges against individually identifiable and reliably measurable risk *components*: a WTI KO call priced with FX held fixed at inception (isolating pure commodity-price risk), paired with an independent Garman—Kohlhagen FX call on a notional fixed at inception (isolating pure currency risk). Two hedging instruments, two hedge relationships.

The distinction is not cosmetic. It is visible directly in the payoff formulas (Section 4): Structure A’s payoff multiplies by the *floating* $S_2(T)$, creating a genuine quanto correlation coupling between the two risk factors inside a single instrument; Structure B’s WTI leg multiplies by the *fixed* $S_2(0)$, deliberately severing that coupling so that the leg prices and hedges commodity risk alone. This paper builds both architectures on an identical, jointly-simulated correlated jump-diffusion path bank, applies the full IFRS 9 cash-flow-hedge accounting mechanics (hypothetical-derivative lower-of testing, cost-of-hedging amortization, knock-out discontinuation, OCI reclassification) to each path at daily granularity, and reports the resulting statement-of-financial-position (SFP), other-comprehensive-income (OCI), ineffectiveness, and economic-residual metrics as measured outcomes.

The underlying quanto KO product (its jump-diffusion dynamics, its Longstaff—Schwartz regression pricing surface, and its double-barrier touch mechanics) is specified fully in Section 3 of this paper; this paper’s entire contribution is on top of that pricing layer, in the accounting-designation question, and Section 3 is written to be self-contained.

The remainder of the paper proceeds as follows. Section 2 places the model in the IFRS 9 and quanto-pricing literature. Section 3 derives the underlying jump-diffusion dynamics, the quanto payoff construction, the double-barrier touch probabilities, and the Longstaff—Schwartz pricing algorithm from first principles. Section 4 specifies the two designation architectures precisely.

Section 5 derives the hypothetical derivative that makes Structure A’s effectiveness test non-tautological. Section 6 specifies the full accounting engine, including its numerical safeguards and the per-step processing sequence. Section 7 defines the economic-residual layer. Section 8 reports measured results. Section 9 and Section 10 discuss sensitivities and the central trade-off. Section 11 outlines the model boundaries, and Section 12 concludes.

2 Literature Review

IFRS 9 (IASB, 2014) requires cash-flow-hedge accounting to recognize the effective portion of a hedging instrument’s fair-value change in other comprehensive income, with any ineffective portion in profit or loss, determined by a lower-of test between the cumulative change in the hedging instrument and the cumulative change in a hypothetical derivative representing the hedged item (IFRS 9.B6.5.5, 9.6.5.11). Aggregated exposures (combinations of a non-derivative exposure and a derivative, or of two risk components managed together) may be designated as a single hedged item under 9.6.3.4; risk components of that same exposure may alternatively be hedged individually under 9.6.3.1, provided each component is separately identifiable and reliably measurable. For option-based hedging instruments, IFRS 9 permits the time-value component to be accounted for as a cost of hedging, amortized to profit or loss on a systematic basis separate from the intrinsic-value-based lower-of test (9.6.5.15—9.6.5.16). On discontinuation of a hedge relationship (including when the hedged forecast transaction is no longer expected to occur, or here, when a knock-out event extinguishes the hedging instrument), amounts accumulated in OCI are addressed under 9.6.5.11(d)/9.6.5.12. This paper implements these mechanics path-by-path at daily granularity and not asserting their qualitative consequences from first principles.

The pricing side draws on three classical results. Merton (1976) supplies the jump-diffusion framework used for WTI, in which discontinuous jump risk renders the market incomplete and a diversifiable jump-risk premium survives the risk-neutral measure change. Longstaff and Schwartz (2001) supply the least-squares Monte Carlo regression used to price and differentiate the barrier-and-early-exercise-aware WTI surface. Garman and Kohlhagen (1983) supply the closed-form currency option formula used for Structure B’s FX leg, which has no barrier feature and therefore needs no regression. Black (1976) supplies the forward-based commodity option convention underlying the WTI premium quoting convention used throughout. The double-KO barrier’s discrete-monitoring treatment follows the continuity correction of Broadie, Glasserman, and Kou (1997), applied on the same daily path bank that determines every accounting statistic reported below, so that the same discretization choice that fixes the option’s price also fixes its ineffectiveness. The quanto correlation coupling between S_1 and S_2 inside Structure A’s single-instrument payoff, the kind of coupling Reiner (1992) prices in closed form for a continuously-monitored, barrier-free quanto claim, is handled here instead by joint Monte Carlo simulation of the correlated pair, since no closed form of that kind survives the addition of a discretely monitored double barrier; the design choice is justified formally in Section 3.

Why a firm designates a hedge relationship at all (accepting the documentation burden, the effectiveness testing, and the cost-of-hedging amortization mechanics of IFRS 9, and not simply

booking the derivative at fair value through profit or loss) is a question this paper’s engine answers quantitatively (by measuring what hedge accounting changes about the P&L path) but does not itself motivate theoretically. That motivation belongs to the corporate hedging literature: Froot, Scharfstein, and Stein (1993) model hedging as a device for protecting a firm’s internal financing capacity, which is the same underlying economic motive that hedge accounting’s smoothing of OCI volatility is designed to make visible in the financial statements and not obscure it; Stulz (1996) develops the practitioner corollary that partial, selective hedging can be value-preserving even when full economic coverage is affordable, directly relevant to Structure B’s asymmetric ineffectiveness split between its WTI and FX legs (Section 8.2). On the accounting-choice question specifically (whether the designation architecture itself, not just the hedge, is shaped by accounting and not purely economic considerations), Glaum and Klöcker (2011) document that hedge-accounting eligibility rules measurably influence which economic hedges firms actually put on, a finding this paper’s own result (two economically similar architectures producing materially different recognized ineffectiveness) is a quantitative, instrument-level instance of.

3 Product and Pricing Definition

This section specifies, completely and from first principles, the quanto double-knock-out product whose designation is the subject of this paper. Every accounting formula in Sections 4—7 is a function of the objects defined here.

3.1 Correlated jump-diffusion dynamics

Let $S_1(t)$ denote the WTI crude price in USD/bbl and $S_2(t)$ the USD/KRW exchange rate, both evolving under the physical measure $\mathbb{P}$. S_1 is a jump-diffusion process; S_2 is a correlated geometric Brownian motion:

$$\frac{dS_1(t)}{S_1(t^-)} = (\mu_1^{\mathbb{P}} - \lambda\kappa)dt + \sigma_1 dW_1(t) + d\left(\sum_{j=1}^{N(t)} (e^{Y_j} - 1)\right) \quad (1)$$

$$\frac{dS_2(t)}{S_2(t)} = \mu_2^{\mathbb{P}} dt + \sigma_2 dW_2(t) \quad (2)$$

$$dW_1(t) \cdot dW_2(t) = \rho dt, \quad \kappa = \mathbb{E}[e^Y - 1] = \exp(\theta_J + \tfrac{1}{2}\delta_J^2) - 1 \quad (3)$$

with $N(t)$ a Poisson process of intensity λ and log-jump-sizes $Y_j \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_J, \delta_J^2)$. The production calibration fits $\lambda = 6.7846$ jumps per year, $\theta_J = -0.0299$, $\delta_J = 0.08443$, a diffusive WTI volatility $\sigma_1 = 0.3242$, an FX volatility $\sigma_2 = 0.0926$, and a WTI—FX correlation $\rho = 0.0876$.

For simulation under the risk-neutral measure $\mathbb{Q}$ used for pricing (drift $\mu_1^{\mathbb{P}} \rightarrow r_{US}$, $\mu_2^{\mathbb{P}} \rightarrow r_{KRW} - r_{US}$), the one-step Euler—Maruyama update with the compound Poisson jump super-

imposed once per step is

$$\varepsilon_1 \sim \mathcal{N}(0, 1), \quad \varepsilon_2^\perp \sim \mathcal{N}(0, 1), \quad \varepsilon_2 = \rho \varepsilon_1 + \sqrt{1 - \rho^2} \varepsilon_2^\perp \quad (4)$$

$$N_i \sim \text{Poisson}(\lambda \Delta t), \quad J_i = \sum_{j=1}^{N_i} Y_j \quad (5)$$

$$S_1(t_{i+1}) = S_1(t_i) \exp \left[\left(r_{US} - \lambda \kappa - \frac{1}{2} \sigma_1^2 \right) \Delta t + \sigma_1 \sqrt{\Delta t} \varepsilon_1 + J_i \right] \quad (6)$$

$$S_2(t_{i+1}) = S_2(t_i) \exp \left[\left(r_{KRW} - r_{US} - \frac{1}{2} \sigma_2^2 \right) \Delta t + \sigma_2 \sqrt{\Delta t} \varepsilon_2 \right] \quad (7)$$

The shared shock ε_1 drives both legs' correlated component; this is the same joint construction that Section 3.3 shows is sufficient, on its own, to price a correlation-coupled quanto payoff without a separate closed-form adjustment.

3.2 Double knock-out barrier

The structure is monitored against an upper barrier $U = 120$ and lower barrier $L = 50$ USD/bbl. Let $\tau_{KO} = \inf\{t : S_1(t) \geq U \text{ or } S_1(t) \leq L\}$. Because the simulation records S_1 only at discrete steps, a path can cross a barrier between two recorded observations without either recorded value lying outside $[L, U]$. The correction applied between each pair of consecutive recorded, interior values S_1^{prev}, S_1 is the exact continuous-Brownian-motion conditional touch probability, in log-space:

$$p_{up} = \exp \left(\frac{-2 \ln(U/S_1^{\text{prev}}) \ln(U/S_1)}{\sigma_1^2 \Delta t} \right) \quad (8)$$

$$p_{dn} = \exp \left(\frac{-2 \ln(S_1^{\text{prev}}/L) \ln(S_1/L)}{\sigma_1^2 \Delta t} \right) \quad (9)$$

Each is compared against an independent uniform draw once per step; a touch is registered if the draw falls below the corresponding probability. Both formulas are exact under continuous Brownian motion between the two endpoints; a genuine Poisson jump landing strictly between two recorded steps can, in principle, cross a barrier without being captured by either the recorded values or the bridge formula. Formally, the true step-conditional touch probability is

$$\mathbb{P}(\tau_{KO} \in (t_i, t_{i+1}] \mid S_1(t_i), S_1(t_{i+1})) = p_{up} + \mathcal{E}_{\text{jump}}(\lambda, \theta_J, \delta_J; \Delta t), \quad (10)$$

where $\mathcal{E}_{\text{jump}} \geq 0$ collects paths a jump crosses the barrier on and then reverts into $[L, U]$ by t_{i+1} , indistinguishable from a jump-free interior path given only the two recorded endpoints. Since such a path needs both a jump landing within an interval of length Δt ($O(\lambda \Delta t)$) and that jump clearing a barrier several diffusive standard deviations away, $\mathcal{E}_{\text{jump}} = O(\Delta t) \rightarrow 0$ as $\Delta t \rightarrow 0$, a small residual gap-risk understatement of the true continuous-time KO frequency. Its size is measured, not merely bounded: a jump-adapted exact re-simulation at 200,000 paths (Poisson jump times drawn exactly within each step, per-sub-interval bridge tests, barrier checked at every jump instant, with the production-style whole-step test evaluated on the *same* paths) puts $\mathcal{E}_{\text{jump}}$ at +0.14 to +0.18 percentage points of KO frequency at this daily grid (risk-neutral and physical banks respectively), confirming the $O(\Delta t)$ analysis with a number (Section 11).

3.3 Quanto payoff construction and the correlation coupling

The importer’s exposure is to the KRW cost of barrels, barrels $\cdot S_1(t) S_2(t)$: both legs enter multiplicatively. The three payoff variants relevant to this paper are (Structure A’s, and the two legs of Structure B’s, are given fully in Section 4); what all of them share is that Structure A’s payoff multiplies the WTI intrinsic value by the *floating* terminal rate $S_2(T)$, so that

$$\mathbb{E}^{\mathbb{Q}}[\max(S_1(T) - K, 0) \cdot S_2(T)] \neq \mathbb{E}^{\mathbb{Q}}[\max(S_1(T) - K, 0)] \cdot \mathbb{E}^{\mathbb{Q}}[S_2(T)] \quad (11)$$

whenever $\rho \neq 0$. A closed-form treatment of this coupling ordinarily requires an explicit quanto drift adjustment (a numeraire change that shifts S_1 ’s effective risk-neutral drift by $-\rho\sigma_1\sigma_2$). This engine never writes that adjustment down explicitly: because S_1 and S_2 are simulated jointly from the correlated shocks of Section 3.1, every path already carries the true correlated joint realization of $(S_1(T), S_2(T))$, and multiplying the two legs together path-by-path before averaging across paths reproduces the correlation-coupled expectation to simulation accuracy automatically. This is why Structure B’s WTI leg, by contrast, multiplies by the *fixed* $S_2(0)$ and not $S_2(T)$: fixing the FX conversion rate at inception deliberately removes the quanto coupling from that leg, so that it is a genuine single-risk-factor commodity hedge, mathematically as well as by designation intent. This payoff-level distinction (floating conversion in Structure A versus fixed conversion in Structure B’s WTI leg) is the precise mathematical counterpart of the IFRS 9 aggregated-exposure-versus-risk-component choice discussed in Section 4.

3.4 Longstaff—Schwartz pricing surface and backward-induction algorithm

The engine prices and differentiates the barrier option by least-squares Monte Carlo. Writing K for the strike ($K = S_1(0)$, at-the-money-forward) and normalizing $v_1 = S_1/K$, $v_2 = S_2/S_2(0)$, the continuation-value surface at each step is a 5-term polynomial,

$$\tilde{C}(S_1, S_2) = \beta_0 + \beta_1 v_1 + \beta_2 v_2 + \beta_3 v_1^2 + \beta_4 v_2^2 + \beta_5 v_1 v_2, \quad (12)$$

fitted by ordinary least squares on in-the-money, not-yet-knocked-out paths, independently at every discrete step. The procedure: (1) simulate N joint paths forward from t_0 to $t_n = T$, recording the first barrier-hitting step via Section 3.2’s bridge test; (2) initialize terminal value at t_n from the payoff; (3) step backward from $i = n - 1$ to 1, regressing the one-step-discounted future value on (v_1, v_2) over the surviving in-the-money subset to obtain $\beta_0(t_i), \dots, \beta_5(t_i)$, and carrying non-fitted paths’ discounted future value backward unchanged; (4) average the step-0 cash flows for the price, retaining every step’s β coefficients as the pricing and hedging surface. Structure A’s option surface and Structure B’s WTI-only surface (fixed- $S_2(0)$ payoff) are fitted *separately*, since their payoffs differ; Structure B’s FX leg needs no regression surface at all, since the Garman—Kohlhagen closed form applies directly to a barrier-free option.

4 Designation Architectures

4.1 Structure A: single quanto KO designation

The hedging instrument is a single European-style quanto KO call. Payoff at maturity, conditional on no knock-out:

$$V_A(T) = \text{barrels} \cdot \max(S_1(T) - K, 0) \cdot S_2(T) \cdot \mathbf{1}\{\tau_{KO} > T\}, \quad K = S_1(0). \quad (13)$$

Fair value along a path is $\max(\tilde{C}(S_1, S_2) \cdot \text{barrels}, \text{intrinsic})$ from Structure A's own regression surface; the floor at intrinsic value is a numerical safeguard described in Section 6.5, where its bias near the upper barrier is also recorded. On knock-out, $V_A \rightarrow 0$; the CFH reserve freezes and any unamortized time value is swept to ineffectiveness (Section 6.3).

4.2 Structure B: dual independent designations

B1 (WTI KO call). Payoff $\text{barrels} \cdot \max(S_1(T) - K, 0) \cdot S_2(0) \cdot \mathbf{1}\{\tau_{KO} > T\}$, priced from a separately-fitted WTI-only regression surface. On knock-out, B1's fair value goes to zero and its remaining time value is swept to ineffectiveness:

$$\text{residual_TV}_W = \max(\text{TV}_{W,0} - \text{COH}_w^{\text{amort}}, 0), \quad (14)$$

using $\text{TV}_{W,0}$, the *inception* time value (initial fair value minus initial intrinsic), not the already-zeroed post-KO fair value.

B2 (Garman—Kohlhagen FX call). Strike $G_0 = S_2(0)e^{(r_{KRW} - r_{US})T_{FX}}$, notional $N_{FX} = \text{barrels} \cdot F_0^{WTI}$ fixed at inception, while the economically relevant dollar exposure $\text{barrels} \cdot S_1(t)$ is stochastic: a deliberate, fixed-notional structural mismatch. Fair value follows the closed-form Garman—Kohlhagen call,

$$V_{B2}(t) = \text{GKCall}(S_2, G_0, \sigma_2, \tau, r_{US}, r_{KRW}, N_{FX}), \quad (15)$$

since the leg carries no barrier and needs no regression surface. B2 survives WTI knock-out; its subsequent fair-value changes flow to FVTPL, not OCI (Section 6.3).

Both B1 and B2 use intrinsic-value-based lower-of comparisons, not a mix of intrinsic value on one side and full fair value (including time value) on the other, which would otherwise bias the two legs' measured effectiveness against each other for reasons unrelated to the underlying hedge quality.

5 The Hypothetical Derivative

IFRS 9's lower-of test compares the cumulative change in the hedging instrument against the cumulative change in a hypothetical derivative (HD) that mirrors the hedged risk without hedge-specific features (IFRS 9.B6.5.5). Using Structure A's own fair value as its HD would make the comparison tautological: ΔV_A against itself is identically equal, so ineffectiveness would be zero

by construction on every path, including knocked-out ones. The hedged item here is a highly probable purchase that does *not* itself knock out; the correct HD is therefore the barrier-free quanto forward on the physical exposure,

$$H^*(t) = \text{barrels} \cdot S_1(t) S_2(t) \exp(\rho \sigma_1 \sigma_2 (T - t)), \quad (16)$$

the closed-form risk-neutral expectation of undiscounted physical settlement at T , projected from t . Because V_A carries the KO indicator and H^* does not, ΔV_A and $h_t \Delta H^*$ (Section 6.4) are not related by construction: on knocked-out paths $V_A \rightarrow 0$ while H^* continues to track $\text{Phys}(t) = \text{barrels} \cdot S_1(t) S_2(t)$, and cumulative ineffectiveness $\Delta V_A - \Delta \text{CFHR}_A$ can be, and empirically is, materially nonzero.

6 Accounting Engine

6.1 Signed lower-of test

At each step,

$$\text{SignedLowerOf}(x, y) = \begin{cases} 0 & \text{if } \text{sign}(x) \neq \text{sign}(y) \text{ and both nonzero} \\ \text{sign}(x) \cdot \min(|x|, |y|) & \text{otherwise.} \end{cases} \quad (17)$$

The sign-consistency branch is not incidental: if the hedging instrument and the hypothetical derivative (or hedged item) have moved in *opposite* directions since inception, the qualifying economic relationship that the hedge relies on has itself broken down for that comparison, and no positive effective portion should be recognized. Structure A applies this to $(\Delta V_A - V_{A,0}, h_t \Delta H^*)$; Structure B applies it separately to B1 (intrinsic WTI change versus barrels $\cdot (S_1 - S_1(0)) S_2(0)$) and B2 (intrinsic FX change versus intrinsic value evaluated on the floating notional barrels $\cdot S_1 e^{r_{US} \tau}$).

6.2 Cost-of-hedging amortization

Inception time value $\text{TV}_0 = V_0 - \text{Intr}_0$ amortizes straight-line over the horizon, $\text{COH}_{\text{step}} = \text{TV}_0 / \text{steps}$. Remaining deferred time value $\max(\text{TV} - \text{COH}^{\text{amort}}, 0)$ sits in OCI alongside the CFH reserve until amortized or swept on discontinuation.

6.3 Knock-out event processing

On the KO step, Structure A sets $V_A = 0$, sweeps remaining time value per Section 4.1, and freezes CFHR/OCI at their KO-moment values. Structure B zeros B1 and applies its own residual-time-value sweep; B2's COH freezes at its KO-moment leftover ($\max((\text{FV}_x - \text{Intr}_x) - \text{COH}_x^{\text{amort}}, 0)$, since B2 economically survives) while its subsequent fair-value changes accrue to FVTPL_B and not OCI. This routes *both* legs' post-event treatment off a single shared discontinuation flag triggered by the WTI barrier; Section 11 discusses the tension this creates with Section 4's framing of B1 and B2 as independent hedge relationships.

6.4 Dynamic hedge ratio and non-circularity

Structure A's hedge ratio h_t is not backed out by dividing V_A by H^* ; it is the analytic sensitivity of the fitted regression surface itself. At step stp , with $v_{1n} = S_1/S_1(0)$, $v_{2n} = S_2/S_2(0)$,

$$\text{rawH} = \frac{\beta_1(stp) + 2\beta_3(stp)v_{1n} + \beta_5(stp)v_{2n}}{S_1(0) \cdot S_2}, \quad h_t = \text{clip}_{[0,1]}(\text{rawH}), \quad (18)$$

$$\Delta\text{CFHR}_A = \text{SignedLowerOf}(\Delta V_A - V_{A,0}, h_t \cdot (H_t^* - H_0^*)). \quad (19)$$

What h_t is, dimensionally. The β coefficients carry units of KRW per unit of the structure, and the denominator $S_1(0)S_2$ is $(\text{USD/bbl}) \times (\text{KRW/USD}) = \text{KRW/bbl}$, so rawH has dimensions of barrels. Multiplied by H^* , which is in KRW, the product in (19) would carry bbl·KRW and could not be compared with ΔV_A in KRW. The clip to $[0, 1]$ is what reconciles this: after clipping, h_t is used as a dimensionless coverage fraction, which is also what IFRS 9 means by a hedge ratio, since the standard defines it as a quantity relationship between hedging instrument and hedged item. The engine therefore computes a quantity that originates as a delta and applies it as a ratio. This is a deliberate choice and not an oversight, but it should be stated plainly, because h_t enters the lower-of test directly and the headline KRW 23.36bn is a function of it. Two consequences follow. Values of rawH outside $[0, 1]$ are truncated rather than rescaled, so the mapping from delta to coverage is not one-to-one at the boundaries; and, as in the delta-hedging companion paper, the lower clip removes the negative sensitivities an up-and-out call genuinely carries near the barrier. A specification that normalized rawH by the structure's own barrel notional instead of truncating it would be dimensionally cleaner, and the difference between the two is not measured here. Because $\beta_0(stp), \dots, \beta_5(stp)$ are the same coefficients used to compute V_A from Section 3.4's surface, h_t is a genuine $\partial V / \partial S_1$ -type sensitivity fixed at the start of the step, independent of that step's realized V_A ; the system would be circular only if ΔCFHR_A were *defined* as $h_t \Delta H^*$ and not *constrained* against it by the lower-of operator above, and it is not.

6.5 Numerical safeguards

The accounting engine enforces the following by design:

- **Intrinsic floor.** The fitted WTI regression value is floored at intrinsic value, $\max(\tilde{C} \cdot \text{barrels}, \text{intrinsic})$, at every step, not only at expiry, preventing local polynomial under-shoot from driving the fitted surface below intrinsic. The safeguard is a numerical one and is described as such here: for a *barrier-free* call the floor is a genuine no-arbitrage bound, but an up-and-out call may legitimately trade below intrinsic, because a spot level deep in the money is also a spot level close to extinction. At $S_1 = 119$ against $U = 120$ and $K = 78.94$ the intrinsic value is large while the contract is nearly certain to die, and the floor overstates the fair value there. How often this matters has been measured rather than left open. On a 200,000-path bank the floor binds on 1,811,721 of 17,213,700 hedged nodes, 10.5% of the total, lifting the fitted value by a mean 4.7% of intrinsic where it binds. The distribution of those nodes is the point: within ten percent of the upper barrier the floor binds on 950,566 of 950,966 nodes, or 99.96%. The safeguard is therefore not

occasional and not incidental. It is active on essentially every node in the region where it is theoretically invalid, and it lifts V_A and V_{B1} there, in the direction of understating measured ineffectiveness on paths approaching knock-out. Section 10.1 states what follows for the reported figures.

- **Unambiguous regression-coefficient extraction.** The fitted regression coefficients are read from a fixed, validated array orientation at every step and not inferred, eliminating a class of silent coefficient misreads that an ambiguous read would otherwise permit.

6.6 Per-step processing sequence

At each step t , the engine processes, in this order: (1) compute $\tau_{WTI} = \max(0, T_{WTI} - t)$, $\tau_{FX} = \max(0, T_{FX} - t)$ and the corresponding expiry flags; (2) evaluate Structure A’s fair value, intrinsic/time-value split, COH, and h_t from the quanto surface; (3) evaluate B2’s fair value unconditionally via the closed-form Garman—Kohlhagen formula (no barrier, no conditional branching required); (4) evaluate B1’s fair value with knock-out checked *before* own-maturity expiry, then own-maturity expiry, then the interior no-arbitrage-floored regression value, in that priority order; (5) apply the lower-of test and freeze logic per Section 6.1—6.3; (6) update the cost-of-hedging balances; (7) accumulate the economic-residual layer (Section 7); (8) roll forward ledger pointers. This ordering (knock-out before own-maturity expiry) resolves the only place where two independent stopping conditions (a shared barrier event and each leg’s own contractual maturity) could otherwise conflict for Structure B’s two legs.

6.7 Maturity treatment

$T_{WTI} = 0.833$ and $T_{FX} = 0.5$ years remain genuinely different in the production configuration; the engine does not force a common horizon. Structure B’s WTI leg is intended to run to T_{WTI} and its FX leg to the shorter T_{FX} , with the FX leg’s fair value frozen at its own-maturity intrinsic value from T_{FX} onward per the source code cited in Section 6.5—6.6. Whether this freeze is observable in the aggregated cross-path output is examined empirically in Section 8 and discussed in Section 11: it is not.

7 Economic Layer

Accounting effectiveness and economic hedge quality are not the same measurement, and this section deliberately keeps them separate.

7.1 Step residuals and telescoping

With $\text{Phys}(t) = \text{barrels} \cdot S_1(t)S_2(t)$ and $\Delta V_B = \Delta V_{B1} + \Delta V_{B2}$,

$$\Delta E_A = \Delta V_A - \Delta \text{Phys}, \quad \Delta E_B = \Delta V_B - \Delta \text{Phys}. \quad (20)$$

Per-path cumulative sums telescope:

$$\sum_{i=1}^n \Delta E_A^{(i)} = (V_A(T) - V_A(0)) - (\text{Phys}(T) - \text{Phys}(0)), \quad (21)$$

and identically for B , a direct consequence of summing a telescoping first difference, which follows by construction. Full-horizon dispersion is measured on the *terminal path cumulant* $E_A^{\text{path}} = \sum_i \Delta E_A^{(i)}$, not on the per-step cross-sectional standard deviation at any single step:

$$\sigma_{\text{econ}}^A = \text{StDev}_{\text{paths}}(E_A^{\text{path}}), \quad \sigma_{\text{econ}}^B = \text{StDev}_{\text{paths}}(E_B^{\text{path}}). \quad (22)$$

The per-step cross-path standard deviation computed at each of the 217 steps is a separate, diagnostic-only series (Figure 2); because (22) aggregates a full-horizon cumulative sum, a single anomalous step’s cross-sectional spread does not dominate it, a property exploited directly in Section 8 to interpret the engine’s numerical-safeguard history.

7.2 Post-KO naked exposure

For $t > \tau_{KO}$ on knocked-out paths,

$$N_A(t) = \text{Phys}(t), \quad N_B(t) = \max(\text{Phys}(t) - V_{B2}(t), 0), \quad (23)$$

reflecting that Structure A loses all hedge coverage on KO while Structure B’s surviving FX leg provides partial offset. VaR₉₉ of the path-wise maximum post-KO naked exposure is reported for each structure.

8 Empirical Results

All results are extracted directly from the production engine’s ledger output after execution, $n = 20,000$ accounting paths, $S_1(0) = 78.94$ USD/bbl, $S_2(0) = 1,540.64$ KRW/USD.

8.1 Inception and pricing-surface validation

One convention is fixed here for the whole paper. Structure A’s premium of KRW 34,264,941,589 is the American-exercise LSMC valuation of the combined quanto knock-out, KRW 17,132.47 per barrel on 2,000,000 barrels, taken from the companion pricing paper. The engine walks the accounting ledger on the same American surface, so where Section 4.1 describes the instrument as European-style it is describing the payoff form and not the exercise convention used in valuation.

Total inception premium outflow is KRW 34,264,941,589, matching Structure A’s own mean inception fair value, a direct internal consistency check that the engine’s premium booking and its own fair-value engine agree at $t = 0$.

The independent reprice uses a reduced sample of $n_{\text{sens}} = 10,000$ LSMC paths and not the production surface; validation points are evaluated at $\tau/T \in \{0.1, 0.5, 0.9\}$ and not including $\tau/T = 0$, since the inception point is a distinct case, separately and directly verified against the

Table 1: Pricing-surface validation: Beta-surface fair value vs. independent reduced-sample LSMC reprice, Structure A and Structure B1.

τ/T	A surface (bn KRW)	A reprice	A % diff	B1 % diff
0.10	38.71	40.55	4.75	1.14
0.50	33.22	33.08	0.42	4.74
0.90	18.26	19.35	5.97	10.01

booked premium (Section 8.1) and not by evaluating the fitted regression surface at a stressed test spot. Agreement is within 5% for Structure A throughout and for Structure B1 up to mid-life, widening to roughly 10% for B1 near expiry, consistent with ordinary finite-sample reprice noise on a shrinking, more path-dependent regression sample and not a systematic bias in either direction.

8.2 H1: cumulative ineffectiveness

Table 2: Headline cross-path comparison, Structure A vs. Structure B.

Metric (bn KRW unless noted)	Structure A	Structure B
Mean cumulative ineffectiveness	23.36	6.40
Full-horizon σ_{econ}	104.52	103.18
Post-KO naked exposure VaR ₉₉	645.44	634.27
Hedge relationships	1	2

Structure A’s mean cumulative |ineffectiveness| is KRW 23.36 billion versus Structure B’s KRW 6.40 billion. The statistic is $\mathbb{E}[|\sum_i \Delta \text{Ineff}^{(i)}|]$, the cross-path mean of the absolute terminal cumulant, not the absolute value of the mean; Structure B’s figure is the signed combination of its two legs, $|-6.57 + 0.18| = 6.40$. Figure 1 plots the cross-path *mean* of the signed cumulant, which is why its Structure A endpoint reads below the tabulated figure: signed cancellation across paths reduces the mean relative to the mean absolute, roughly three-and-a-half times larger. With the tautology-avoiding HD of Section 5 in place, this does not arise from a degenerate zero-ineffectiveness setup: it reflects that H^* continues tracking the physical exposure through knock-out while V_A is extinguished, so every KO path contributes a large, non-cancelling ineffectiveness term to Structure A that Structure B, whose two legs are tested separately and whose FX leg survives KO, does not generate to the same degree (Figure 1).

8.3 H2: full-horizon economic residual

Full-horizon $\sigma_{\text{econ}}^A = \text{KRW } 104.52$ billion and $\sigma_{\text{econ}}^B = \text{KRW } 103.18$ billion, a difference of roughly 1.3%, well within ordinary Monte Carlo re-seeding noise observed across repeated engine runs (independent re-executions of the identical engine produced σ_{econ}^B values ranging from KRW 102.4 to KRW 103.2 billion). The two architectures are, empirically, economically indistinguishable at the full-horizon aggregate level despite the several-fold difference in accounting




Figure 1: Mean cumulative P&L ineffectiveness over time, Structure A vs. Structure B (B1+B2).

ineffectiveness reported in Section 8.2: direct evidence that accounting effectiveness and economic residual risk measure genuinely different quantities. Because the two architectures run on a shared path bank, the comparison can be made per path rather than across runs, and it has been. On 200,000 paired paths the validated re-implementation returns $\sigma_{\text{econ}}^A = \text{KRW } 107.01\text{bn}$ against $\sigma_{\text{econ}}^B = \text{KRW } 104.40\text{bn}$, and the paired difference is unambiguous: $\mathbb{E}[|E_A| - |E_B|] = \text{KRW } 9.76\text{bn}$ with a standard error of $\text{KRW } 0.09\text{bn}$, a t -statistic of 103.9, and a paired test of equal variance rejecting at $t = 145.0$. The two architectures do *not* carry statistically equal economic risk, and any statement to that effect would be an artifact of comparing cross-run re-seeding dispersion instead of the paired difference the shared bank makes available.

What survives, and what the paper’s argument actually needs, is the economic magnitude rather than the statistical one. Structure A’s residual risk exceeds Structure B’s by 2.5% of level on the re-implementation and by 1.3% on the production ledger. Set against a recognized-ineffectiveness difference of several hundred percent between the same two architectures, a two-percent difference in residual risk is not what the accounting is responding to. The claim made throughout this paper is therefore the ordering-of-magnitudes one, a claim about economic size rather than about statistical indistinguishability.

Figure 2 shows a smooth per-step series throughout the horizon for both structures, with no discontinuity at the terminal step. This is a useful property of the metric design in Section 7.1: because full-horizon σ_{econ} (Eq. 22) aggregates a 217-step cumulative sum per path before taking the cross-path standard deviation, a single step’s cross-sectional spread (even a large one, were one to occur at a boundary condition such as the final reporting step) could not by itself dominate the full-horizon figure; the per-step series is diagnostic precisely because it is more sensitive to isolated numerical artifacts than the aggregate measure that is actually reported as σ_{econ} .

8.4 H3: post-KO naked exposure and reclassification

Knock-out frequency on the shared path bank is 45.1%. The accounting ledger is walked on the *physical*-measure path bank, since the effectiveness test is applied to realized outcomes




Figure 2: Per-step cross-path standard deviation of ΔE (diagnostic; not the full-horizon σ_{econ} of Eq. 22).

rather than to risk-neutral expectations, so this figure is the $\mathbb{P}$ -measure barrier-touch rate and is comparable to the 45.12% reported for the $\mathbb{P}$ -bank exact engine in the companion delta-hedging paper, not to that paper’s risk-neutral 43.69%. Post-KO VaR₉₉ of maximum naked exposure is KRW 645.44 billion for Structure A versus KRW 634.27 billion for Structure B: a modest, not dramatic, advantage for Structure B, consistent with the FX leg covering only the currency component of the naked exposure while commodity-price risk remains fully unhedged for both structures once WTI has knocked out. On KO, Structure A’s SFP derivative line goes to zero and KRW 24.41 billion of accumulated OCI is reclassified into profit or loss, a cross-path mean over the $n = 20,000$ simulated paths. Appendix B documents the relationship between this reclassification amount and the engine’s per-step SFP identity.

The cross-path means above compress every path’s individual knock-out event into a single number; Figure 3 shows one representative path instead, chosen to knock out near mid-horizon. Structure A’s derivative and Structure B1’s WTI leg, sharing the same barrier event, both cliff to zero at the knock-out step, the hedge relationship extinguished on both designations simultaneously. Structure B2, the FX leg, carries no barrier of its own and survives: it continues marking to fair value via the closed-form Garman—Kohlhagen formula for the remainder of the path, its subsequent changes now flowing straight to profit or loss (FVTPL) and not OCI, since the shared discontinuation flag has already fired. This is the concrete, single-path mechanism behind the aggregate post-KO FVTPL statistic of Section 8.5 (mean KRW 0.28 billion, standard deviation KRW 5.14 billion on knocked-out paths): a designation architecture that survives the barrier event keeps generating P&L volatility precisely when the firm has lost its primary coverage, while one that extinguishes cleanly does not.

8.5 H4: SFP and disclosure footprint

Structure A carries one derivative line and one hedge relationship; Structure B carries two, one of which (B2) remains live and volatile on the SFP after WTI knock-out. To ensure robust estimation of the post-KO FVTPL metrics, this paper reports a jump-adapted exact




Figure 3: One representative knock-out path. Top: WTI spot against the double barrier. Bottom: fair value of Structure A, Structure B1 (WTI leg), and Structure B2 (FX leg). A and B1 cliff to zero at the shared knock-out event; B2 survives and continues marking to fair value, now outside hedge accounting.

re-derivation at $n = 200,000$ paths, using the same numerical safeguards validated elsewhere in this paper, yielding a mean post-KO FVTPL of KRW 0.28 billion and a standard deviation of KRW 5.14 billion on knocked-out paths.




Figure 4: Mean SFP derivative carrying amount: Structure A (single line) vs. Structure B (B1+B2, shown both summed and separately).




Figure 5: Mean OCI reserve balance over time, Structure A vs. Structure B.

8.6 The maturity-freeze mechanism

Section 6.7 describes a per-leg maturity-freeze mechanism: once t passes $T_{FX} = 0.5$, B2's fair value freezes at its own-maturity intrinsic value while B1 continues to $T_{WTI} = 0.833$. Figure 7 plots B1 and B2's mean fair value across the full horizon.

The freeze is unconditional: from step 131 (the first step at or beyond T_{FX}) onward, B2's cross-path mean fair value is constant at every subsequent step, KRW 8.638 billion, for ev-




Figure 6: Structure A pricing-surface validation error by elapsed maturity fraction.




Figure 7: B1 (WTI) and B2 (FX) mean fair value over time. B2 freezes at $t = T_{FX} = 0.5$, independent of each path's WTI knock-out status.

ery path regardless of whether that path’s WTI leg has separately been knocked out. Freeze completion is gated on FX-leg maturity alone and not on WTI knock-out survival, so paths already knocked out before T_{FX} freeze their B2 mark at T_{FX} just as surviving paths do, and the cross-path mean shows the sharp plateau visible in Figure 7: every path contributes the identical frozen value from T_{FX} onward, so the average is that value, with no smoothing from asynchronous freeze timing across paths.

This same unconditional gating has a direct consequence for the post-KO FVTPL statistic of Section 8.5. That statistic pools two populations the freeze mechanism makes structurally distinct. For a path with $\tau_{KO} < T_{FX}$, B2 is still marking to fair value when the WTI barrier fires, is reclassified to FVTPL while still live, and continues generating genuine FVTPL volatility until it separately freezes at T_{FX} , the “zombie window” Section 8.4 describes. For a path with $\tau_{KO} \geq T_{FX}$, B2 has *already* frozen, on its own contractual schedule, before the WTI barrier ever fires; its fair value is unchanged through and past the knock-out event, so $FVTPL_B$ accrues no further profit or loss for the remainder of that path, by the same arithmetic that produces the plateau in Figure 7: reclassification to FVTPL happens on paper, but there is no leg left alive to generate any P&L from it. The 200,000-path exact re-derivation of Section 8.5 finds close to an even division of the knocked-out population (52.6% knock out before T_{FX} , 47.4% after); and, as the freeze mechanism predicts, zero post-KO FVTPL mean and zero standard deviation on every path in the $\tau_{KO} \geq T_{FX}$ subgroup, with the entire distribution’s dispersion concentrated in the $\tau_{KO} < T_{FX}$ subgroup (mean KRW 0.53 billion, standard deviation KRW 7.08 billion on that subgroup alone). The pooled headline is therefore not a description of what a “typical” knocked-out path experiences; it is a mixture of a subpopulation that experiences the full zombie-FVTPL mechanism and a same-sized subpopulation that experiences none of it, already inert by the time its designation status changes: both the population split and the zero result follow from the freeze mechanism itself.

9 Sensitivity

Reprice sample size. The independent LSMC reprice sample used for Table 1 is $n_{\text{sens}} = 10,000$ paths. The A % diff column is computed from the two level columns beside it, so the table is internally consistent; an earlier draft carried a % column from a separate reprice run, which is why Figure 6 and this column now agree only up to reprice noise. The τ/T index is elapsed time fraction, so $\tau/T = 0.9$ is near expiry. Near-expiry validation error is single-digit percent for Structure A throughout, and for Structure B1 up to mid-life, widening only modestly near expiry, consistent with ordinary Monte Carlo reprice noise near a payoff kink at this sample size and not a systematic bias in either direction.

Static vs. dynamic hedge ratio. Fixing $h_t \equiv h_0$ at its inception value would remove the regression surface’s sensitivity to moneyness ($2\beta_3 v_1$) and cross-asset correlation ($\beta_5 v_2$) from Structure A’s effectiveness test; the production engine recomputes h_t every step from Section 6.4, and a full static-ratio counterfactual run is left for future work. This is the single most useful robustness check not performed here: because h_t enters (16)’s comparison directly, the headline KRW 23.36bn is a function of the dynamic ratio, and a static- h_0 run would separate

the mismatch attributable to the payoff’s multiplicative form from the mismatch attributable to re-estimating the ratio each step.

9.1 A validated fourth cell: the combination drives the gap

Structure A is always a single instrument *and* always barriered; Structure B is always split *and* its WTI leg (B1) is also always barriered. The main comparison therefore never isolates whether Structure A’s elevated ineffectiveness (Section 8.2) comes from combining the exposure into one instrument, from the knock-out feature that instrument carries, or from both together. Separating a 2×2 design (combined-vs-split, barriered-vs-vanilla) requires two further structures beyond A and B: a combined quanto call with the barrier removed, and a split WTI-only call (FX fixed, B1’s payoff) with the barrier removed.

Decomposing Structure B’s own ledger first, and not only citing its total, is the right place to start: terminal cumulative ineffectiveness splits KRW -6.57 billion for B1 against KRW $+0.18$ billion for B2 (frozen at T_{FX} , Section 8.6): essentially all of Structure B’s reported KRW 6.40 billion comes from B1, not B2. B1 carries the identical knock-out barrier as Structure A, yet its ineffectiveness is barely a quarter of Structure A’s (6.57bn against 23.36bn), a $3.56\times$ gap between two barriered instruments differing only in whether FX is fixed or floating.

Resolving that gap requires an engine faithful enough to reproduce it before it can be trusted to answer the harder question: any re-implementation that shows no combined-versus-split effect at all, when a $3.56\times$ gap demonstrably exists in production, has omitted something material and should not be relied upon for the cells production cannot itself compute. Extracting the production engine’s mechanics line-by-line identifies the three features such fidelity requires: the hypothetical derivative H^* is a *closed form*, barrels $\cdot S_1 \cdot S_2 \cdot e^{\rho\sigma_1\sigma_2\tau}$, never a regression fit; Structure A’s lower-of test compares *intrinsic*-value changes (not fair-value changes) against $h_t \cdot (H_t^* - H_0^*)$, the hypothetical scaled by the dynamic hedge ratio of Section 6.4; and at the knock-out step the remaining unamortized cost-of-hedging is swept directly into that step’s ineffectiveness. A re-implementation faithful to all three mechanics, at daily (217-step) monitoring, in a two-stage design matching production (fit Beta_Mat on one path bank, walk a separate fresh accounting bank forward against the frozen surface), reproduces every available production statistic to within 1—16%: Structure A ineffectiveness 24.78bn (production 23.36bn), Structure B1 ineffectiveness -7.70 bn (production -6.57 bn), economic residual std 107.3bn/104.4bn (production 104.52bn/103.18bn), and critically, an A-to-B1 ratio of $3.22\times$ against production’s $3.56\times$. Passing this validation check is what licenses trusting its answer to the fourth cell.

	KO / barriered	Vanilla / barrier-free
Combined (A-analog)	24.78bn	55.77bn
Split (B1-analog)	7.70bn	≈ 0.00 bn

Table 3: Mean |terminal cumulative ineffectiveness| (bn KRW), production-fidelity re-implementation (daily monitoring, exact h_t /intrinsic/closed-form- H^* mechanics), all four cells on one validated shared engine. Compare to production: Structure A (combined, barriered) KRW 23.36bn; Structure B1 (split, barriered) KRW 6.57bn.

The result reverses the natural first guess entirely. **Split + vanilla is indistinguishable**

from zero. B1’s lower-of test compares option intrinsic value, $\text{barrels} \cdot \max(S_1 - K, 0) \cdot S_2(0)$, against a fixed-notional WTI forward, $\text{barrels} \cdot (S_1 - S_1(0)) \cdot S_2(0)$. These are the identical function of S_1 whenever the option is in the money, and the sign-consistency rule of Section 6.1 resolves the out-of-the-money case to zero as well: a single-asset option’s intrinsic value cannot help but track a forward on the same asset, with or without a barrier. **Combined stays large in both barrier states.** Structure A’s hedged item is the *product* of two floating quantities, $\max(S_1 - K, 0) \cdot S_2$; no forward-style hypothetical, scaled or not, is that same function of (S_1, S_2) , so the mismatch never mechanically vanishes the way B1’s does, in either barrier state.

The mismatch can in fact be named exactly, which is worth doing because it converts an empirical regularity into an analytic one. On paths that are alive and in the money, Structure A’s intrinsic value decomposes as

$$\begin{aligned} \text{barrels} \cdot (S_1 - K) S_2 &= \underbrace{\text{barrels} \cdot S_1 S_2}_{= H^* \text{ by (33)}} - \text{barrels} \cdot K S_2, \end{aligned} \quad (24)$$

so the hedged item and the hypothetical derivative differ by precisely the term $-\text{barrels} \cdot K \cdot S_2(t)$: the currency exposure of the *strike* leg. Structure B1 fixes the conversion at $S_2(0)$, which makes that term a constant and therefore cancels it out of every first difference; Structure A leaves it floating, so it survives into ΔV_A and has no counterpart in ΔH^* . This is the mechanism behind the $3.2\times$ gap, stated more sharply than “because the payoff is a product”.

The magnitude checks out. The strike leg carries $\text{barrels} \cdot K = \$157,880,000$, worth KRW 243.2bn at inception, and $\sigma_2 \sqrt{T_{WTI}} = 8.45\%$, giving a one-standard-deviation currency swing on that leg of about KRW 20.6bn against the measured KRW 23.36bn. The correspondence is not an exact attribution, since $h_t < 1$ scales the comparison and the lower-of test truncates it, but the order and the sign are those of (24), and no other term in the decomposition is of that size. The combined-vanilla cell’s *absolute* level, however, is the one cell with no production anchor and proves sensitive to re-implementation details (fit safeguards, path law, sample size move it across a wide range while every other cell’s ordering is stable), so this paper draws no conclusion about whether removing the barrier raises or lowers the combined structure’s ineffectiveness: only that it never approaches the split structure’s mechanical zero.

Structure A’s ineffectiveness is therefore driven by the combination of two floating exposures in one multiplicative payoff, and not by the barrier. For the split structure the relationship is reversed: B1’s ineffectiveness is negligible in the absence of a barrier and is created almost entirely *by* the barrier’s knock-out-day cost-of-hedging sweep. Economic risk responds to the barrier in the opposite direction from accounting ineffectiveness: combined+vanilla’s residual-risk standard deviation (42.3bn) is far below combined+KO’s (107.3bn), since removing the barrier removes the total-loss-of-coverage-on-knock-out mechanism of Section 8.3. Whether a structure is barriered and whether it is combined therefore pull accounting ineffectiveness and economic risk in opposite directions, and Section 10 returns to what this means for reading the two structures’ numbers together.

10 Discussion

The accounting architecture of this paper rests on a single balance, $V(t) = \text{OCI}(t) + \text{RE}(t)$, which Appendix B confirms at every step and which holds by construction of the pro-forma rather than as an independent test of the engine. The check with content is the flow-level decomposition, which fails once knock-outs begin by an amount equal to the KRW 24.41 billion reclassification: Appendix B reports that gap as a finding. OCI is where the effective, well-behaved part of a hedge’s fair-value movement is parked (a reserve, not an income statement line) alongside the deferred cost-of-hedging (Section 6.2) amortizing quietly in the background. P&L is where the lower-of test’s *ineffective* residual is required to land the moment it is measured, whether or not the firm has done anything wrong: it records how exactly a specific hypothetical derivative happens to track a specific instrument, and is better read as a measurement of that tracking convention than as a report of hedging skill. Retained earnings is where both amortized cost-of-hedging and every recognized ineffectiveness charge finally accumulate: the one line that actually marks the firm’s capital. The two architectures compared here do not disagree about total economic risk; full-horizon σ_{econ} differs by 1.3% of level between them, KRW 104.52bn against 103.18bn, a paired difference that is statistically sharp but economically small (Section 8.3). They disagree, sharply, about *when* and *how visibly* that risk is forced out of OCI’s holding pattern and into P&L and, from there, into retained earnings.

Structure A’s OCI reserve behaves normally for as long as the option survives: intrinsic-value changes track the barrier-free quanto forward reasonably well, cost-of-hedging amortizes on schedule, and the running ineffectiveness charge is modest. The break comes in one step, at the knock-out event. V_A collapses to zero; the CFHR reserve freezes at its last value; and (a mechanic easy to miss in the ledger but decisive in the total) whatever cost-of-hedging has not yet amortized is swept directly into that same step’s ineffectiveness and not carried forward (Section 6.3). Section 9.1’s validated re-derivation locates *why* this charge is so much larger than Structure B’s: the driver is that Structure A’s hedged item is the product of two floating quantities, $\max(S_1 - K, 0) \cdot S_2$, a function that no forward-style hypothetical derivative can replicate under any fixed scaling. Removing the barrier does not repair this mismatch: the split leg’s zero is unreachable for a product payoff in either barrier state. The barrier’s actual role is to concentrate the unavoidable mismatch into one visible, dated event. What does not vanish at that event is the frozen OCI reserve accumulated up to it (KRW 24.41 billion, Section 8.4): IFRS 9’s own mechanics direct it toward the acquired inventory’s cost basis at physical settlement and not toward an immediate further P&L charge, so Structure A’s capital-account damage is concentrated, but bounded, at a moment the firm can see coming as the barrier approaches.

Structure B’s day-to-day ledger looks better than Structure A’s for a more mechanical reason than the intuition that two smaller hedges are safer than one. B1’s own lower-of test compares option intrinsic value against a fixed-notional WTI forward (the same function of S_1 whenever the option is in the money), so the two track each other exactly by construction; the validated re-derivation puts B1’s ineffectiveness in the absence of a barrier at a number indistinguishable from zero, with essentially all of its production-observed KRW 6.57 billion coming from the barrier’s own knock-out-day cost-of-hedging sweep, the same mechanic as Structure A’s, just

smaller because B1 has no floating-FX term to fight. This clean-looking book is the point at which a CFO reading only the ineffectiveness line could be misled: it says nothing about B2. B2 carries no barrier, is unaffected by the WTI knock-out on its own terms, and continues marking to fair value via Garman—Kohlhagen for as long as it is designated, but two independent events end that designation early and separately. First, at its own contractual maturity $T_{FX} = 0.5$, well before $T_{WTI} = 0.833$, B2’s hedge relationship is discontinued by maturity, not by any market event (Section 8.6). Second, and more consequentially for paths that survive past T_{FX} , the moment the WTI barrier fires, B2 is reclassified out of hedge accounting entirely and into FVTPL (Section 8.5): not because anything happened to the FX rate, but because its *partner* leg’s discontinuation flag fired. From that point its fair-value changes flow straight into P&L with no OCI buffer at all, and this is where Structure B’s real capital-account risk actually lives: post-knock-out FVTPL volatility with a standard deviation of KRW 5.14 billion pooled across all knocked-out paths, rising to KRW 7.08 billion on the roughly half of paths for which the zombie window is genuinely open (Section 8.5), silently eroding retained earnings for the remainder of the horizon, well after the CFO’s attention has moved past the quiet ineffectiveness line that looked so favorable earlier.

A stronger reading of the same evidence is worth stating explicitly, because it changes what the ineffectiveness numbers mean. The gap between Structure A’s and Structure B’s measured ineffectiveness is not a difference in economic risk. Section 8.3 finds full-horizon σ_{econ} differing by only 1.3% of level across the two architectures, against an accounting gap of several hundred percent, and Section 9.1 traces the ineffectiveness gap not to any feature of the underlying exposure but to the algebraic form of the hedged item relative to the hypothetical derivative the standard permits: a product of two floating quantities cannot be tracked by a forward-style hypothetical, while a single floating quantity can. Both structures hedge the same physical cargo, at the same notional, against the same two market factors. What differs is the designation boundary the firm draws around them and, consequently, which hypothetical derivative IFRS 9 obliges it to construct.

The hypothetical derivative is a stipulated benchmark and not a traded instrument, defined by the standard’s own conventions (perfect terms-match, no counterparty credit, a specified treatment of time value and cost of hedging), and the “ineffectiveness” the lower-of test recognizes is the distance between an actual instrument and that stipulated object. That distance is a property of the measurement convention at least as much as of the hedge. On this reading, Structure A’s larger ineffectiveness is better described as an accounting artifact than as a discovered inefficiency: the multiplicative payoff behaves no worse economically than the split legs, while being less representable within the specific comparison the standard mandates. The result is a divergence between economic substance, which the two structures share, and accounting recognition, which separates them by a factor of several.

The causality therefore runs the other way than a direct reading of the accounting output would suggest. A firm that splits a designation has selected an architecture whose hedged item is representable by an admissible hypothetical derivative; its underlying risk is unchanged. The measured ineffectiveness ordering is real, auditable, and consequential for reported earnings, and for those reasons it is a genuine constraint on structuring decisions; but it is a constraint

imposed by the reporting framework and not by the market. The practical implication is that the ineffectiveness line should be read as a designation-architecture diagnostic and not as a risk metric, and that any firm choosing between these structures on the basis of reported ineffectiveness alone is optimizing against the measurement convention and not against its exposure, which is precisely the behavior Section 8.5’s post-knock-out FVTPL result shows can be costly.

There is also a price on the choice, and it has gone unreported. At inception Structure A costs KRW 34.26bn, the premium of the single combined quanto. Structure B costs KRW 35.62bn for the WTI knock-out call plus KRW 6.44bn for the Garman–Kohlhagen FX call, KRW 42.05bn in total: the split architecture is KRW 7.79bn, or 22.7%, more expensive. Set beside Section 8.3’s finding that the two carry economic risk differing by 1.3% of level, the practical statement is sharper than the ineffectiveness comparison alone suggests. A firm choosing Structure B is paying roughly KRW 8bn for a reduction in recognized ineffectiveness that Section 10 argues is an artifact of the hypothetical-derivative convention rather than a reduction in risk, and is buying, along with it, the post-knock-out FVTPL exposure of Section 8.5 that Structure A does not carry. Stated that way the split designation is an expensive way to buy a quieter ineffectiveness line.

The choice between these architectures is therefore not a choice about how much risk the importer carries (both carry statistically the same economic risk), but about the shape of its accounting recognition. Structure A reports a large, one-time, dated charge to P&L at the barrier, after which its books go quiet: a single visible scar. Structure B reports a small, reassuring ineffectiveness charge for most of the horizon, then hands an unbounded, undated, back-loaded P&L bleed to whichever period the WTI barrier happens to fire in, through a leg that was never the one being watched. A single combined designation trades a bigger number for certainty about when it lands; two independent designations trade a smaller number for a tail risk that surfaces once the firm’s attention, and its designation architecture, has already moved on.

10.1 What the floor and the clip do to the reported numbers

Two engine conventions, the intrinsic floor of Section 6.5 and the $[0, 1]$ clip on h_t of Section 6.4, act in the same region and in the same direction, and their measured extent bounds how much of the headline they can carry.

The floor binds on 99.96% of hedged nodes within ten percent of the upper barrier and lifts fitted value by a mean 4.7% of intrinsic where it binds. Because Structure A’s lower-of test compares ΔV_A against ΔH^* , and the floor raises V_A on precisely the paths where the option is closest to extinction, the floor compresses the measured mismatch on those paths. Its removal would raise, not lower, Structure A’s recognized ineffectiveness; the reported KRW 23.36bn is therefore a lower bound with respect to this convention, and the several-fold A-over-B ordering that the paper rests on is not at risk from it.

The clip on h_t works the other way but is smaller. The companion delta-hedging paper measures the same regression surface on the same calibration and finds the raw gradient never negative over 17.3 million nodes, so the lower clip never binds; the upper clip binds on 1.4% of

nodes. Re-running the delta hedge with the clip removed moves the hedge-cost standard deviation by four hundredths of one percent. On this evidence the clip is not a material determinant of the ledger, and the floor is the convention that warrants the caveat.

What neither measurement can do is establish the level of V_A independently, because the same fitted surface supplies both the value and, through h_t , the ratio. The companion paper’s finding that the surface gradient overstates $\partial V/\partial S_1$ by roughly an order of magnitude therefore transmits directly into h_t here, and is the reason Section 11 keeps the static-ratio counterfactual as the outstanding robustness check rather than treating h_t as settled.

11 Limitations

Several boundaries define this paper’s scope. Monte Carlo estimates (VaR₉₉ at $n = 20,000$; near-expiry reprice validation at $n_{\text{sens}} = 10,000$) carry ordinary sampling noise, so the Structure A/B naked-exposure VaR₉₉ gap ($\sim 2\%$) and the B1 near-expiry reprice discrepancy ($\sim 10\%$) are order-of-magnitude comparisons and not point estimates. The KRW 24.41 billion reclassification figure of Section 8.4 is a cross-path mean; the amount on any individual path varies with knock-out timing (Appendix B). The Brownian-bridge barrier correction (Section 3.2) is exact between recorded steps and understates true knock-out frequency by a measured $+0.14$ to $+0.18$ percentage points at the daily grid, confirmed by a jump-adapted exact re-simulation at 200,000 paths; the same exact engine directionally re-confirms the combined-versus-split ineffectiveness ordering that drives Section 8.2’s headline result, so this paper’s qualitative conclusions do not rest on the production scheme’s discretization. WTI and FX maturities ($T_{WTI} = 0.833$, $T_{FX} = 0.5$) are reported as designed, reflecting the structures’ actual contractual terms and not a counterfactual unified maturity. The 2×2 benchmark of Section 9.1 and the post-KO FVTPL statistic of Section 8.5 run on a re-implementation that reproduces the production engine’s mechanics and validates to within 1–16% of every production statistic the current engine can still generate; the post-KO FVTPL cells and the barrier-free vanilla cells have no production analog to validate against by construction, so this paper reports its own measurement as the primary figure for those cells and not as a check against a production number. The representative knock-out path of Figure 3 illustrates the qualitative mechanism of Sections 8.4–8.5 on a single reconstructed path and is not a source of any reported cross-path statistic. Two engine conventions are quantified in Section 10.1 rather than left open. The intrinsic floor is a valid no-arbitrage bound only for a barrier-free call, and it binds on 99.96% of hedged nodes within ten percent of the upper barrier, lifting fitted value by a mean 4.7% of intrinsic there; because it raises V_A on the paths closest to extinction it compresses the measured mismatch, so the reported KRW 23.36bn is a lower bound with respect to it and the A-over-B ordering is not at risk. The $[0, 1]$ clip on h_t is immaterial by measurement: the raw gradient never turns negative, the lower clip never binds, and removing the clip in the companion paper’s hedge moves the cost standard deviation by 0.04%. What remains genuinely open is the static-ratio counterfactual, since h_t inherits the surface-gradient level error the companion paper documents. American early exercise, basis spreads, and counterparty credit risk are outside this paper’s scope. The paper’s central findings are stable across these boundaries: the combined designation’s mean

ineffectiveness is several times larger than the split designation’s, the gap traces to the combined structure’s multiplicative payoff and not to the barrier, and full-horizon economic risk differs between the two architectures by 1.3% of level, a paired difference that is precisely estimated and economically small, against an accounting difference of several hundred percent. A final boundary is interpretive and not numerical. Every ineffectiveness figure reported here is defined relative to IFRS 9’s hypothetical-derivative construct, and is therefore conditional on that construct’s conventions and not on the economics of the hedge; a different admissible hypothetical, or a jurisdiction permitting a different effectiveness-assessment method, would produce different numbers from identical cash flows. The comparison between architectures is meaningful within the framework, and the paper makes no claim that the ineffectiveness ordering corresponds to an ordering of economic risk. Section 10 argues, on the contrary, that it does not.

12 Conclusion

This paper specifies and measures, and not assumes, the accounting consequences of two IFRS 9 cash-flow-hedge designation architectures for the same aggregated commodity-and-currency exposure. The hypothetical derivative of Section 5 removes Structure A’s tautological zero-ineffectiveness outcome; the dynamic hedge ratio of Section 6.4 is non-circular by construction; and the full-horizon economic-residual measure of Section 7.1 separates accounting effectiveness from economic risk explicitly. On $n = 20,000$ simulated paths, Structure A shows roughly three-and-a-half times larger mean cumulative ineffectiveness than Structure B, while the two structures are economically indistinguishable on full-horizon residual risk and only modestly different on post-KO naked-exposure tail risk. Decomposing Structure B further attributes essentially all of that smaller ineffectiveness to its WTI leg and not its FX leg, and a production-fidelity re-implementation validated to within 1—16% of every reproducible production statistic traces the remaining gap to its source: combining two floating exposures into one multiplicative quanto payoff, not the knock-out barrier, is what defeats any forward-style hypothetical derivative’s ability to track Structure A, since a single-asset structure’s intrinsic value tracks a forward almost by definition while a product of two floating prices cannot. The barrier’s role is to concentrate this unavoidable mismatch into one dated, visible event and not to cause it. The choice between a single combined designation and two independent risk-component designations is, on this evidence, a choice about *when* and how visibly recognized ineffectiveness reaches the firm’s capital accounts (a front-loaded, bounded shock at a known event for Structure A, versus a quiet ledger followed by an undated, back-loaded FVTPL bleed through an unwatched leg for Structure B), not a choice about how much economic protection the importer ultimately retains.

A Mathematical Derivations

A.1 Derivation of the hypothetical-derivative closed form

Section 5 states $H^*(t)$ in closed form; this subsection derives it from the same $\mathbb{Q}$ -measure dynamics used throughout Section 3, so that the formula rests on the stated SDEs and not on

assertion.

Write $\tau = T - t$ and let $R_1 = S_1(T)/S_1(t)$, $R_2 = S_2(T)/S_2(t)$ denote the two legs' relative returns over $[t, T]$ under $\mathbb{Q}$. Integrating Section 3.1's correlated jump-diffusion log-dynamics over $[t, T]$,

$$R_1 = \exp \left[(r_{US} - \lambda\kappa - \tfrac{1}{2}\sigma_1^2)\tau + \sigma_1\Delta W_1 + J \right], \quad J = \sum_{j:t < T_j \leq T} Y_j, \quad (25)$$

$$R_2 = \exp \left[(r_{KRW} - r_{US} - \tfrac{1}{2}\sigma_2^2)\tau + \sigma_2\Delta W_2 \right], \quad (26)$$

with $\Delta W_1 = W_1(T) - W_1(t)$, $\Delta W_2 = W_2(T) - W_2(t)$ jointly normal, each with variance τ and covariance $\rho\tau$, and J , the compound-Poisson jump sum accumulated over the interval, independent of the Brownian increments. Since H^* requires $\mathbb{E}^\mathbb{Q}[S_1(T)S_2(T) \mid \mathcal{F}_t] = S_1(t)S_2(t) \mathbb{E}^\mathbb{Q}[R_1R_2 \mid \mathcal{F}_t]$, the deterministic exponent and the stochastic factor separate:

$$\mathbb{E}^\mathbb{Q}[R_1R_2] = \exp \left[(r_{KRW} - \lambda\kappa - \tfrac{\sigma_1^2}{2} - \tfrac{\sigma_2^2}{2})\tau \right] \cdot \mathbb{E}^\mathbb{Q}[e^{\sigma_1\Delta W_1 + \sigma_2\Delta W_2}] \cdot \mathbb{E}^\mathbb{Q}[e^J]. \quad (27)$$

Both remaining expectations are standard closed forms. $\sigma_1\Delta W_1 + \sigma_2\Delta W_2$ is normal with mean 0 and variance $(\sigma_1^2 + \sigma_2^2 + 2\rho\sigma_1\sigma_2)\tau$, so its exponential moment is $\exp[\frac{1}{2}(\sigma_1^2 + \sigma_2^2 + 2\rho\sigma_1\sigma_2)\tau]$; and J is a compound-Poisson sum at rate $\lambda\tau$ with $\mathbb{E}[e^Y] = 1 + \kappa$ by the definition of κ in Section 3.1, giving the standard compound-Poisson exponential moment $\mathbb{E}[e^J] = \exp(\lambda\kappa\tau)$. Substituting and collecting terms,

$$\mathbb{E}^\mathbb{Q}[R_1R_2] = \exp \left[(r_{KRW} - \lambda\kappa - \tfrac{\sigma_1^2}{2} - \tfrac{\sigma_2^2}{2})\tau + (\tfrac{\sigma_1^2}{2} + \tfrac{\sigma_2^2}{2} + \rho\sigma_1\sigma_2)\tau + \lambda\kappa\tau \right] \quad (28)$$

$$= \exp \left[(r_{KRW} + \rho\sigma_1\sigma_2)\tau \right], \quad (29)$$

where the jump-compensator terms $-\lambda\kappa\tau$ and $+\lambda\kappa\tau$ cancel exactly, as do the two half-variance terms $\pm\frac{1}{2}\sigma_1^2\tau$ and $\pm\frac{1}{2}\sigma_2^2\tau$, leaving only the correlation cross-term $\rho\sigma_1\sigma_2\tau$, the quanto adjustment, produced here as a byproduct of the joint expectation and not by an explicit drift shift under a changed numeraire. (29) is confirmed numerically against 4×10^6 direct Monte Carlo paths at the production calibration ($\lambda = 6.7846$, $\theta_J = -0.0299$, $\delta_J = 0.08443$, $\sigma_1 = 0.3242$, $\sigma_2 = 0.0926$, $\rho = 0.0876$, $\tau = 0.833$): closed-form 1.031842 against simulated 1.031760, a relative gap of 8×10^{-5} , within Monte Carlo noise at this sample size. Hence

$$\mathbb{E}^\mathbb{Q}[S_1(T)S_2(T) \mid \mathcal{F}_t] = S_1(t)S_2(t) \exp \left[(r_{KRW} + \rho\sigma_1\sigma_2)\tau \right]. \quad (30)$$

Discounting this KRW-denominated expectation back to t at the KRW risk-free rate (the correct rate, since barrels $\cdot S_1(T)S_2(T)$ is itself a KRW cash amount) multiplies by $e^{-r_{KRW}\tau}$, which cancels the $r_{KRW}\tau$ term in the exponent exactly, leaving (16):

$$H^*(t) = e^{-r_{KRW}\tau} \cdot \text{barrels} \cdot \mathbb{E}^\mathbb{Q}[S_1(T)S_2(T) \mid \mathcal{F}_t] = \text{barrels} \cdot S_1(t)S_2(t) \exp(\rho\sigma_1\sigma_2\tau). \quad (31)$$

The domestic risk-free rate appears nowhere in the final expression: it enters once as S_2 's $\mathbb{Q}$ -drift and again, with opposite sign, as the discount rate, and the two cancel identically regardless of the rate's level. What survives is a term $\rho\sigma_1\sigma_2\tau$, and the paper's earlier reading of it requires

correction. It is tempting to treat that term as evidence that joint simulation on correlated paths reproduces the quanto drift adjustment automatically. It is in fact the trace of that adjustment being *absent*. Since the dynamics of Section 3.1 give S_2 the drift $r_{KRW} - r_{US}$, the simulation measure is $\mathbb{Q}^{KRW}$; under it $S_1 S_2$ is a KRW-denominated traded asset and must drift at r_{KRW} , which forces

$$\text{drift}(S_1) = r_{KRW} - (r_{KRW} - r_{US}) - \rho\sigma_1\sigma_2 = r_{US} - \rho\sigma_1\sigma_2. \quad (32)$$

The engine uses r_{US} without the $-\rho\sigma_1\sigma_2$ term, and (29) carries exactly that omission forward. Correlated shocks fix the joint distribution’s shape, not its location; a drift adjustment is an artifact of the change of measure and has to be imposed rather than inherited.

With (32) in force the algebra simplifies rather than complicates: $\mathbb{E}[R_1 R_2] = e^{r_{KRW}\tau}$, the discount cancels it entirely, and the hypothetical derivative collapses to

$$H^*(t) = \text{barrels} \cdot S_1(t) S_2(t), \quad (33)$$

with no correlation term at all. Numerically the correction is $e^{\rho\sigma_1\sigma_2\tau} - 1 = 0.22\%$ of H^* . Its effect on the reported results has been measured rather than assumed: re-running the validated re-implementation on 200,000 paths with the correlation factor removed from H^* , as (32) requires, moves Structure A’s mean absolute ineffectiveness by $+0.76\%$. No ordering, sign, or headline comparison in this paper changes. The correction matters for what the appendix is entitled to claim: (29) validates the algebra of the joint expectation, and it does not validate the measure specification. Joint simulation on correlated paths does not, on its own, reproduce the quanto adjustment. This is also what makes H^* fit for purpose under IFRS 9.6.3.4: the hypothetical derivative must benchmark the aggregated exposure’s multiplicative WTI-times-FX risk linearly and without hedge-specific optionality, and a discounted $\mathbb{Q}$ -expectation of the raw physical settlement is, by construction, a $\mathbb{Q}$ -martingale carrying no barrier or other option-specific feature: the property Section 5 requires of it, derived here and not assumed.

B SFP Ledger Balance Verification

A natural conjecture is that the derivative asset (debit), cost-of-hedging amortization and any P&L loss (expense) sum to the OCI reserve (credit) and any P&L gain (revenue) at *every* step: a full double-entry decomposition of the accounting flows themselves, not merely of the resulting balances. Tested directly against the production ledger, that specific decomposition does not hold once knock-out events begin occurring across paths: the gap between the two sides grows from zero in the pre-knock-out period to a terminal value that equals, to the reported digit, the KRW 24.41 billion reclassification amount of Section 8.4. That gap is a quantifiable signature of the residual-time-value knock-out sweep of Section 4.2 moving value from an implicit running balance into an explicit one-time reclassification event, and is reported here as a finding in its own right and not as the flow-level identity originally conjectured.

The identity that *does* hold, verified directly against the ledger at every one of the 217 steps for both structures, is simpler and is enforced by construction in the SFP pro-forma output:

the derivative asset’s carrying amount equals the OCI reserve plus retained earnings,

$$V_A(t) = \text{OCI}_A(t) + \text{RE}_A(t), \quad V_{B1}(t) + V_{B2}(t) = \text{OCI}_B(t) + \text{RE}_B(t), \quad (34)$$

with retained earnings computed directly as $\text{RE} = V - \text{OCI}$ and not accumulated independently from signed P&L line items. Measured against the production ledger the two sides agree to within 10^{-6} — 10^{-8} KRW absolute at every tested step for both structures, against carrying amounts on the order of 10^{10} — 10^{11} KRW.

That agreement should be read for exactly what it is, and no more. Because RE is *defined* as $V - \text{OCI}$ in the SFP pro-forma, (34) is true by construction and the 10^{-6} residual measures floating-point arithmetic rather than the accounting engine. It is a useful guard against transcription and aggregation faults in the pro-forma, and it is not an independent check on the ledger’s economics: no sign-convention error in the P&L attribution could cause it to fail. The check with genuine content is the flow-level decomposition in the paragraph above, and that one does *not* hold once knock-outs begin. This paper therefore reports the flow-level gap as its substantive ledger finding, and treats (34) as a consistency guard on the output format.




Figure 8: SFP debit (derivative asset, black outline) vs. credit (OCI reserve plus retained earnings, stacked area) over time, Structure A and Structure B. The outline traces the top of the stacked area at every step by construction (Eq. 34). Structure B’s visible kink near $t \approx 0.5$ is the maturity-freeze plateau of Section 8.6.

Figure 8 plots both sides of (34) for each structure; the black outline (derivative asset) traces the top of the stacked OCI-plus-retained-earnings area, and Structure B’s stacked area visibly flattens at the FX leg’s own maturity, the same mechanism documented in Section 8.6. Figure 9 restates the identity for Structure A at maturity as a single-step flow diagram: the terminal derivative asset of KRW 12.79 billion splits into an OCI reserve of KRW 24.41 billion and a cumulative retained-earnings position of −KRW 11.62 billion, which sum, with sign, back to the derivative asset.




Figure 9: Structure A terminal split of the derivative asset into OCI reserve and cumulative retained earnings, verified via Eq. 34.

References.

- Black, F. (1976). The pricing of commodity contracts. *Journal of Financial Economics*, 3(1—2), 167—179.
- Broadie, M., Glasserman, P., and Kou, S. G. (1997). A continuity correction for discrete barrier options. *Mathematical Finance*, 7(4), 325—349.
- Froot, K. A., Scharfstein, D. S., and Stein, J. C. (1993). Risk management: coordinating corporate investment and financing policies. *Journal of Finance*, 48(5), 1629—1658.
- Garman, M. B., and Kohlhagen, S. W. (1983). Foreign currency option values. *Journal of International Money and Finance*, 2(3), 231—237.
- Glaum, M., and Klöcker, A. (2011). Hedge accounting and its influence on financial hedging: when the tail wags the dog. *Accounting and Business Research*, 41(5), 459—489.
- IASB (2014). *IFRS 9 Financial Instruments*. London: IFRS Foundation.
- Longstaff, F. A., and Schwartz, E. S. (2001). Valuing American options by simulation: a simple least-squares approach. *Review of Financial Studies*, 14(1), 113—147.
- Merton, R. C. (1976). Option pricing when underlying stock returns are discontinuous. *Journal of Financial Economics*, 3(1—2), 125—144.
- Reiner, E. (1992). Quanto mechanics. *Risk*, 5(3), 59—63.
- Stulz, R. M. (1996). Rethinking risk management. *Journal of Applied Corporate Finance*, 9(3), 8—25.